



(Formerly ITM University, Gurugram)

Database Management

System

Lab Manual

Database

Department of Computer Science and Engineering

The NorthCap University, Gurugram

**Management System Lab Manual
CSL 214**

2019-20

Ms. Kanika Gupta
Dr. Prachi
Dr. Anuradha Dhull
Dr. Priyanka Vashisht



(Formerly ITM University, Gurugram)

Department of Computer Science and Engineering
The NorthCap University, Gurugram- 122017, India
Session 2019-20

Published by:

School of Engineering and Technology
Department of Computer Science & Engineering
The NorthCap University Gurugram

- Laboratory Manual is for Internal Circulation only**

© Copyright Reserved

No part of this Practical Record Book may be reproduced, used, stored without prior permission of The NorthCap University

Copying or facilitating copying of lab work comes under cheating and is considered as use of unfair means. Students indulging in copying or facilitating copying shall be awarded zero marks for that particular experiment. Frequent cases of copying may lead to disciplinary action. Attendance in lab classes is mandatory.

Labs are open up to 7 PM upon request. Students are encouraged to make full use of labs beyond normal lab hours.

2019-20

PREFACE

Database Management System Lab Manual is designed to meet the course and program requirements of NCU curriculum for B.Tech II year students of CSE branch. The concept of the lab work is to give practical experience for basic lab skills to students. It provides the space and scope for self-study so that students can come up with new and creative ideas.

The Lab manual is written on the basis of “teach yourself pattern” and expected that students who come with proper preparation should be able to perform the experiments without any difficulty. Brief introduction to each experiment with information about self-study material is provided. The laboratory exercises will include familiarization with SQL and NoSQL commands to define, manipulate, retrieve and manage data in a database. The course will enable students to create a database and devise queries for extracting information from the database using Relational Algebra and SQL. The students will be able to appreciate the importance of handling unstructured data using NoSQL and apply the concepts of DBMS for developing a backend for a non-trivial project.

At the start of each experiment a question bank for preparation and practice is suggested which may be used to test the basic understanding of the students about the experiment. Students are expected to come thoroughly prepared for the lab. General disciplines, safety guidelines and report writing are also discussed.

The lab manual is a part of curriculum for the TheNorthCap University, Gurugram. Teacher's copy of the experimental results and answer for the questions are available as sample guidelines.

We hope that lab manual would be useful to students of CSE, IT, ECE and BSc branches and author requests the readers to kindly forward their suggestions / constructive criticism for further improvement of the work book.

Author expresses deep gratitude to Members, Governing Body-NCU for encouragement and motivation.

Authors
The NorthCap University
Gurugram, India

CONTENTS

S.No.	Details	Page No.
	Syllabus	
1	Introduction	
2	Lab Requirements	
3	General Instructions	
4	List of Experiments	
5	Rubrics	
6	Annexure 1 (Format of Lab Report)	
7	Annexure 2 (Format of Lab Certificate)	

SYLLABUS

1. Department:	Department of Computer Science and Engineering		
2. Course Name: Database Management Systems		3. Course Code	4. L-T-P
		CSL214	3-0-2
6. Type of Course (Check one):	Programme Core <input checked="" type="checkbox"/>	Programme Elective <input type="checkbox"/>	Open Elective <input type="checkbox"/>
7. Pre-requisite(s), if any: None			
8. Frequency of offering (check one): Odd <input type="checkbox"/> Even <input checked="" type="checkbox"/> Either semester <input type="checkbox"/> Every semester <input type="checkbox"/>			
9. Brief Syllabus:			
<p>Databases form the core of all major applications – finance, social, administrative, education etc. Organizations work on large volumes of data every day, introducing the need to have database management systems to easily identify, extract, store and transform details in the database. This course will explore concepts and principles of DBMS, database design, data modeling, database implementation, and database management through various assignments and projects. By the end of this course, the student will be able to work as a database engineer by designing, developing and maintaining the database for any project application.</p>			
Total lecture, Tutorial and Practical Hours for this course (Take 15 teaching weeks per semester): 75			
Lectures: 45hours		Practice	
		Tutorials: 10 hours	Lab Work: 20 hours
10. Course Outcomes (COs)			
<p>Possible usefulness of this course after its completion i.e. how this course will be practically useful to him once it is completed.</p>			
CO 1	Identifying contrast between traditional and modern Database Systems, thereby recognizing their applications.		
CO 2	Developing conceptual database design for any real time project by defining the relationship, constraints etc. on entities.		
CO 3	Applying appropriate design techniques to design a good database that meets the user requirement.		
CO 4	Creating a database and devising queries for extracting information from the database using Relational Algebra and SQL.		

CO 5	Applying the concepts of DBMS for developing a backend for a non-trivial project using NoSQL.	
CO 6	Ability to improvise data fetching time by applying indexing concepts.	
CO 7	Understanding the concepts of end-to-end transaction processing in a database.	

11. UNIT WISE DETAILS No. of Units: 7

Unit Number: 1	Title: Introduction to Database Systems	No. of hours: 4
-----------------------	--	------------------------

Content Summary:

Overview of Database Management Systems, Advantages of DBMS over File Processing Systems, DBMS Vs. RDBMS, DBA roles and responsibilities, Data Independence, Architecture of Database(3-Schema Architecture, Complete architecture), Database Query Languages (DDL, DML, DCL), Relational Model Concepts: Primary Key, Unique key, Foreign key, Super Key, Alternate key, Candidate key, Constraints used in Relational Data Model including integrity constraints.

Unit Number: 2	Title: Conceptual Database Design	No. of hours: 8
-----------------------	--	------------------------

Content Summary:

Data Modeling Using the Entity Relationship (ER) Model, The Enhanced Entity-Relationship (EER) Model: Entity Set, attributes and their types, Relationship Constraints (including Participation constraints and cardinality ratio), ER Diagrams, constraints and design issues, Reduction of ER and EER diagram to relational schemas.

Unit Number: 3	Title: Relational Database Design	No. of hours: 8
-----------------------	--	------------------------

Content Summary:

Relational database design, Functional dependencies: Fully functional dependency, partial FD, trivial, non-trivial FD, inference rules, canonical cover, lossless join, dependency preservation, multi value dependency, Normal Forms: 1NF, 2NF, 3NF, BCNF, 4NF, 5NF, Normalization and denormalization process

Unit Number: 4	Title: Query Languages	No. of hours: 8
-----------------------	-------------------------------	------------------------

Content Summary:

Relational Algebra: relational operators (projection, selection , union, intersection, set difference, division, rename, Cartesian product, generalized relational algebra operators, cross product, join operators : inner vs. outer join, theta join, outer join, natural join, equijoin, self-join, complete set of relational algebra operations).

SQL- Queries, Constraints, Form of SQL query, UNION, INTERSECT and EXCEPT, Nested queries, Aggregate Operators, Null values, Complex Integrity constraints in SQL and triggers.

Unit Number: 5	Title: Introduction to NoSQL (MongoDB)	No. of hours: 4
-----------------------	---	------------------------

Content Summary:

Introduction to MongoDB, Datatypes, Document Data Model-Creating, Inserting, Updating and Deleting Documents, MongoDB Query Language, Sorting, Join Operations.

Unit Number: 6	Title: File Organization& Indexing	No. of hours: 6
-----------------------	---	------------------------

2019-20

Content Summary:

Disk Storage, Basic File Structures and Hashing: Unordered, ordered and hashed files of records, Single and multilevel indexes: primary index, secondary index, clustered, multilevel and dynamic multilevel indexes (B-Tree and B+ Tree).

Unit Number: 7	Title: Transaction Management & Concurrency Control	No. of hours: 7
----------------	---	-----------------

Content Summary:

Introduction to transaction processing, ACID Properties, Concurrency control mechanisms: serializability, two phase locking protocol, basic concept of deadlock, deadlock handling, timestamp-based protocols, precedence graph to ensure serializability, different protocols in concurrency control. Database back-up and Recovery.

12. Brief Description of Self-learning components by students (through books/resource material etc.):

Aggregation and Pagination in MongoDB

13. Books Recommended:

Textbooks:

1. Elmasri R. and Navathe S.B., Fundamentals of Database Management Systems. 6th ed. Pearson, 2010.
2. Silberschatz A., Korth H.F. and Sudarshan S., Database System Concepts. 6th ed. Mc.Graw Hill, 2010.
3. Chodorow K., MongoDB: The Definitive Guide. 2nd ed. O'Reilly Media, 2013.

Reference Books:

1. Ramakrishnan R. and Gehrke J., Database Management Systems. 3rd ed. McGraw-Hill Education, 2003.
2. Suehring S., My SQL Bible. Wiley Publishing, 2002.

Reference Websites: (nptel, swayam, coursera, edx, udemy, lms, official documentation weblink)

- <https://nptel.ac.in/courses/106105175/2>
- <https://docs.mongodb.com/>

2019-20

1. INTRODUCTION

That ‘learning is a continuous process’ cannot be over emphasized. The theoretical knowledge gained during lecture sessions need to be strengthened through practical experimentation. Thus, practical makes an integral part of a learning process.

The purpose of conducting experiments can be stated as follows:

- To familiarize the students with the basic concepts, programming skill development and the take home laboratory assignments mainly implementation-oriented which have to be coded in high level language. The lab sessions will be based on exploring the concepts discussed in class.
- Developing conceptual database design for any real time project by defining the relationship, constraints etc. on entities.
- Compare and contrast application of SQL and NoSQL for managing a database.
- Hands on experience on the experimental setup and software.

2. LAB REQUIREMENTS

S.No.	Requirements	Details
1	Software Requirements	MySQL, MongoDB
2	Operating System	Any
3	Hardware Requirements	Windows and Linux: Intel 64/32 or AMD Athlon 64/32, or AMD Opteron processor 2 GB RAM 80 GB hard disk space
4	Required Bandwidth	NA

3. GENERAL INSTRUCTIONS

3.1 General discipline in the lab

- Students must turn up in time and contact concerned faculty for the experiment they are supposed to perform.
- Students will not be allowed to enter late in the lab.
- Students will not leave the class till the period is over.
- Students should come prepared for their experiment.
- Experimental results should be entered in the lab report format and certified/signed by concerned faculty/ lab Instructor.
- Students must get the connection of the hardware setup verified before switching on the power supply.
- Students should maintain silence while performing the experiments. If any necessity arises for discussion amongst them, they should discuss with a very low pitch without disturbing the adjacent groups.
- Violating the above code of conduct may attract disciplinary action.
- Damaging lab equipment or removing any component from the lab may invite penalties and strict disciplinary action.

3.2 Attendance

- Attendance in the lab class is compulsory.
- Students should not attend a different lab group/section other than the one assigned at the beginning of the session.

2019-20

- On account of illness or some family problems, if a student misses his/her lab classes, he/she may be assigned a different group to make up the losses in consultation with the concerned faculty / lab instructor. Or he/she may work in the lab during spare/extra hours to complete the experiment. No attendance will be granted for such case.

3.3 Preparation and Performance

- Students should come to the lab thoroughly prepared on the experiments they are assigned to perform on that day. Brief introduction to each experiment with information about self study reference is provided on LMS.
- Students must bring the lab report during each practical class with written records of the last experiments performed complete in all respect.
- Each student is required to write a complete report of the experiment he has performed and bring to lab class for evaluation in the next working lab. Sufficient space in work book is provided for independent writing of theory, observation, calculation and conclusion.
- Students should follow the Zero tolerance policy for copying / plagiarism. Zero marks will be awarded if found copied. If caught further, it will lead to disciplinary action.
- Refer **Annexure 1** for Lab Report Format

4. LIST OF EXPERIMENTS

S.No.	Division of Experiments	List of Experiments
1	Conceptual Database Design	Design an ER diagram for the COMPANY database for the following set of requirements.
2		Design a Relational Database Design for the COMPANY database from the ER/EER diagram for the following set of requirements.
3	SQL	To apply SQL integrity constraints as per the DDL statements given below for COMPANY database.
4		To familiarize with SELECT-FROM-WHERE SQL simple queries on the COMPANY database.

2019-20

5		To familiarize with different JOIN operations in SQL on the COMPANY database.
6		To understand Aggregate functions and Group by Clause using SQL queries on the COMPANY database.
7		To familiarize with nested SQL queries on the COMPANY database.
8		Identifying contrast between Relational Databases and NoSQL, thereby recognizing their applications.
9	NoSQL	Create COMPANY database using NoSQL database - MongoDB.
10		Retrieve data from NoSQL database - MongoDB.

5. RUBRICS

Marks Distribution	
Continuous Evaluation(50 Marks)	End Semester Exam (20 Marks)
Each experiment shall be evaluated for 10 marks and at the end of the semester proportional marks shall be awarded out of 50.	
<p>Following is the breakup of 10 marks for each</p> <p>4 Marks: Observation & conduct of experiment. Teacher may ask questions about experiment.</p> <p>3 Marks: For report writing</p> <p>3 Marks: For the 15 minutes quiz to be conducted in every lab.</p>	<p>End semester practical evaluation including Mini project (if any) carries 20 marks.</p>

2019-20

Annexure 1

Database Management System (CSL 214)

Lab Practical Report



Faculty Name

Kanika Gupta

Student Name Chirag Sardana

Roll No.: 19CSU071

Semester: 4th

Group: FS A1

Department of Computer Science and Engineering
The NorthCap University, Gurugram- 122017, India
Session 2019-20

2019-20

INDEX

Database Management System Lab Manual

2019-20

2019-20

Experiment No: 1

Student Name and Roll Number: Chirag Sardana and 19CSU071

Semester /Section: 4th/FS A1

Link to Code: <https://github.com/chiragsardana/Question-In-Java/tree/master/DBMS/Experiment%201>

Date:

Faculty Signature:

Remarks:

Objective

Design an ER diagram for the COMPANY database for the following set of requirements.

Program Outcome

- The students will be able to draw conceptual database design using ERD Plus.

Problem Statement

The COMPANY database keeps track of a company's employees, departments, and projects. Suppose that after the requirements collection and analysis phase, the database designers provide the following description of the miniworld—the part of the company that will be represented in the database.

- A department controls a number of projects, each of which has a unique name, a unique number, and a single location.
- The company is organized into departments. Each department has a unique name, a unique number, and a particular employee who manages the department. We keep track of the start date when that employee began managing the department. A department may have several locations.
- We store each employee's name, Social Security number, address, salary, sex (gender), and birth date. An employee is assigned to one department, but may work on several projects, which are not necessarily controlled by the same department. We keep track of the current number of hours per week that an employee works on each project. We also keep track of the direct supervisor of each employee (who is another employee).
- We want to keep track of the dependents of each employee for insurance purposes. We keep each dependent's first name, sex, birth date, and relationship to the employee.

Design the ER model by identifying the following from the above requirements:

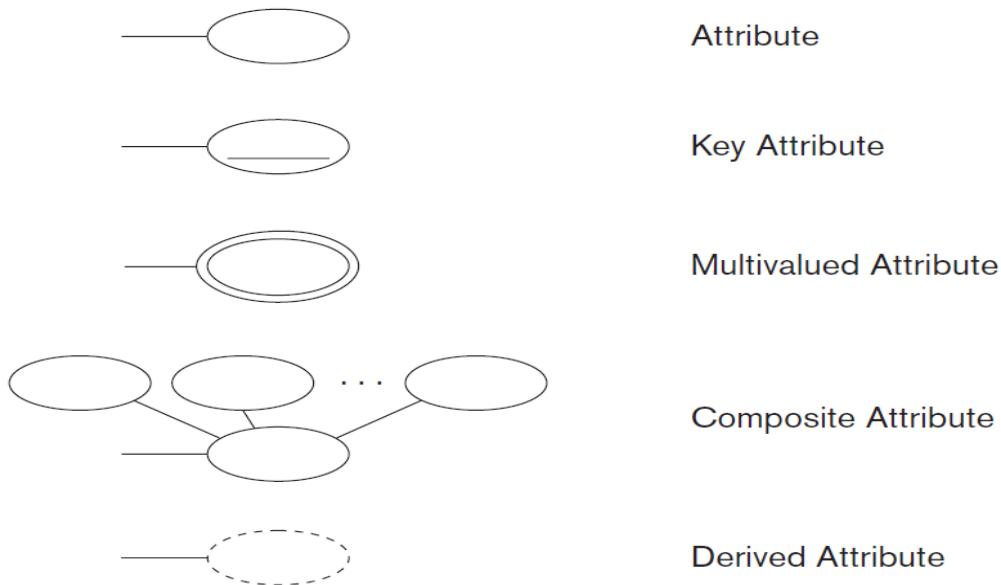
- i) Entities (Strong and Weak)
- ii) Relationships
- iii) Participation constraints
- iv) Various types of attribute
- v) Recursive relations
- vi) Mapping cardinalities

- vii) Binary/Ternary relationship
- viii) Specialization/Generalization etc.

Background Study

ER Diagram Symbols and Notations:

- 1) Entity:
 - Real-world object distinguishable from other objects.
 - An entity is described using a set of *attributes*.
- 2) Entity Set: A collection of similar entities. Eg: all employees.
 - All entities in an entity set have the same set of attributes.
 - Each entity set has a *key*.
 - Each attribute has a *domain*.
- 3) Attributes
 - Attributes are properties used to describe an entity.
 - Example: EMPLOYEE entity may have a Name, SSN, Address, Sex, BirthDate.

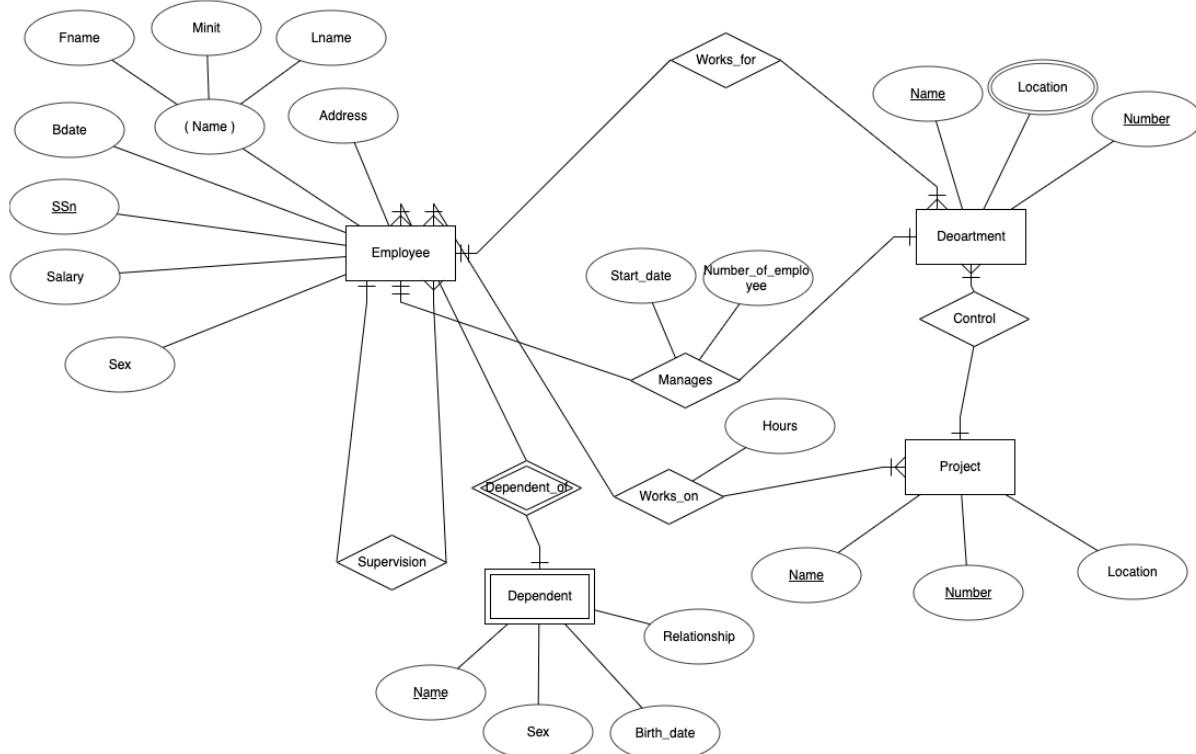


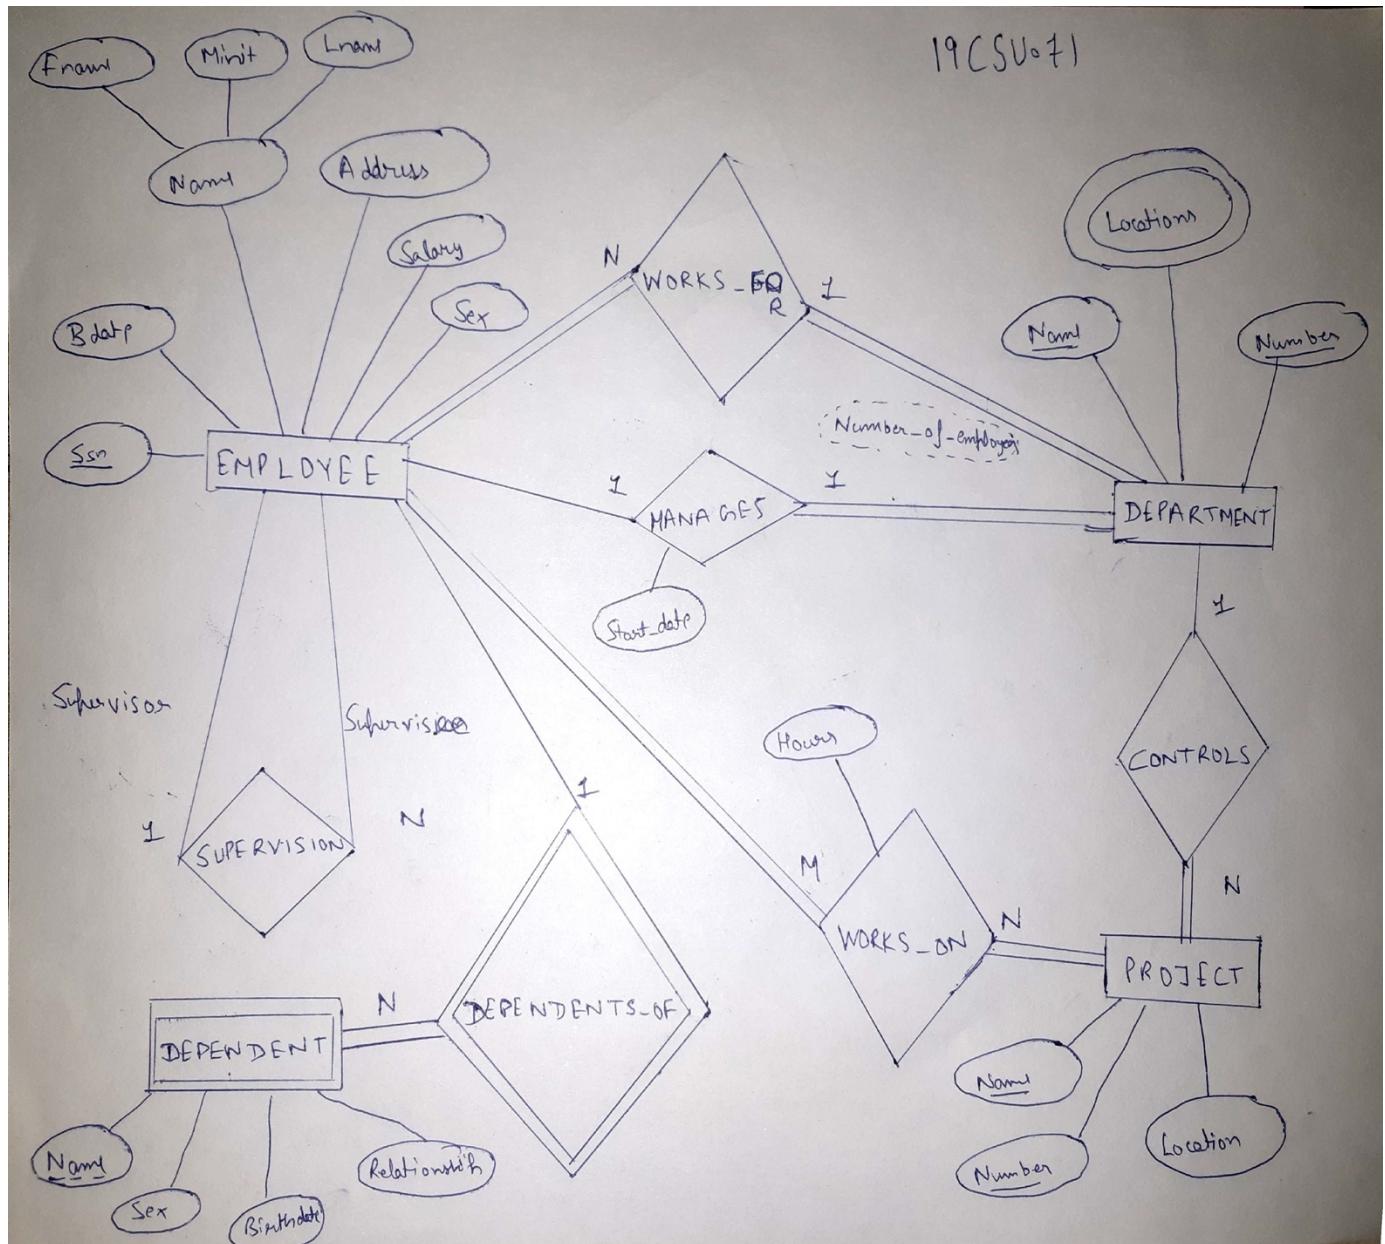
- 4) Relationship
 - A relationship relates two or more distinct entities with a specific meaning.
 - Relationships of the same type are grouped or typed into a relationship type.
 - 3 types of relationships : Unary, Binary & Ternary.
- 5) Recursive Relationship
 - A relationship with the same participating entity type in distinct roles.
 - Example: the SUPERVISION relationship
- 6) Structural Constraints – Semantics of Relationships

2019-20

- Cardinality Ratio : The number of instances of an entity from a relation that can be associated with the relation.
 - Participation Constraints
 - Total Participation – Each entity is involved in the relationship.
 - Partial participation – Not all entities are involved in the relationship.
- 7) Specialization - is the process of defining a set of subclasses of a superclass
- 8) Generalization
 - the reverse of the specialization process .
 - Several classes with common features are generalized into a superclass (original classes become its subclasses)

ER/EER Diagram





Relationships:

Manages	1:1 Employee and Department
Works_for	N:1 Employee and Department
Controls	1:N Department and Project
Supervision	1:N Recursive (Employee)
Works_on	M:N Employee and Project
Dependents_of	1:N Employee and Dependent

2019-20

Preparatory Questions

Q1) Given the basic ER and relational models, which of the following is INCORRECT?

- 1) An attribute of an entity can have more than one value
- 2) An attribute of an entity can be composite
- 3) In a row of a relational table, an attribute can have more than one value
- 4) In a row of a relational table, an attribute can have exactly one value or a NULL value

Answer is : 3

The term ‘entity’ belongs to ER model and the term ‘relational table’ belongs to relational model.

1 and 2 both are true. ER model supports both multivalued and composite attributes See this for more details.

3 is false and 4 is true. In Relation model, an entry in relational table can have exactly one value or a NULL.

Q2) Consider a directed line(-->) from the relationship set advisor to both entity sets instructor and student. This indicates _____ cardinality

- 1) One to many
- 2) One to one
- 3) Many to many
- 4) Many to one

Answer is : 2

Q3) An entity set that does not have sufficient attributes to form a primary key is termed as :

- 1) Strong entity set
- 2) Variant set
- 3) Weak entity set
- 4) Variable set

Answer is : 3

Q4) In which of the following cases an entity instance must be a member of only one subtype?

- 1) Disjoint with total specialisation
- 2) Disjoint with partial specialisation
- 3) Overlap with total specialisation
- 4) Overlap with Partial Specialisation

Answer is : 1

Q5) Which of the following indicates the maximum number of entities can be involved in a relationship?

- 1) Minimum cardinality
- 2) Maximum Cardinality
- 3) ERD
- 4) Greater Entity Count (GEC)

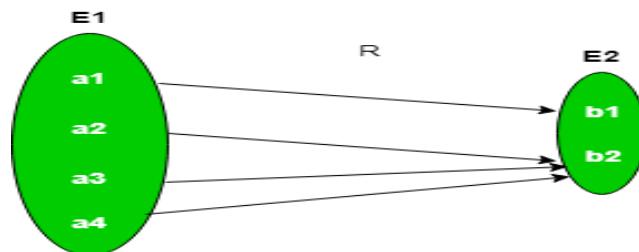
Answer is : 2

Q6) In an Entity-Relationship (ER) model, suppose R is a many-to-one relationship from entity set E1 to entity set E2. Assume that E1 and E2 participate totally in R and that the cardinality of E1 is greater than the cardinality of E2. Which one of the following is true about R?

- 1) Every entity in E1 is associated with exactly one entity in E2.
- 2) Some entity in E1 is associated with more than one entity in E2.
- 3) Every entity in E2 is associated with exactly one entity in E1.
- 4) Every entity in E2 is associated with at most one entity in E1.

Answer is: 1

Since given relation is **many to one** :



Therefore, no entity in E1 can be related to more than one entity in E2 and an entity in E2 can be related to more than one entity in E1.

Only option 1 is correct.

Q7) Which symbol denote derived attributes in ER Model?

- 1) Double ellipse
- 2) Dashed ellipse
- 3) Squared ellipse
- 4) Ellipse with attribute name underlined

Answer is : 2

Q8) How are roles specified in an ER diagram

- 1) By labelling the rectangles
- 2) By labelling the diamonds
- 3) Roles cannot be specified in an ER diagram
- 4) By labelling the lines

Answer is: 4

E-R diagrams can include role indicators. Roles can be specified in an ER diagram by labelling the lines.

Q9) For a weak entity set to be meaningful, it must be associated with another entity set in combination with some of their attribute values, is called as:

- 1) Neighbour Set
- 2) Strong Entity Set
- 3) Owner Entity Set
- 4) Weak Set

Answer is: 3

Q10) Which of the following statements is FALSE about weak entity set?

- 1) Weak entities can be deleted automatically when their strong entity is deleted.
- 2) Weak entity set avoids the data duplication and consequent possible inconsistencies caused by duplicating the key of the strong entity.
- 3) A weak entity set has no primary keys unless attributes of the strong entity set on which it depends are included
- 4) Tuples in a weak entity set are not partitioned according to their relationship with tuples in a strong entity set

Answer is: 4

- Weak entities can be deleted automatically when their strong entity is deleted. **Correct**
- Weak entity set avoids the data duplication and consequent possible inconsistencies caused by duplicating the key of the strong entity. **Correct**
- A weak entity set has no primary keys unless attributes of the strong entity set on which it depends are included **Correct**
- Tuples in a weak entity set are not partitioned according to their relationship with tuples in a strong entity set. This is **Incorrect**, because tuples in a weak entity set are partitioned according to their relationship with tuples in a strong entity set.

2019-20

Experiment No: 2

Student Name and Roll Number: Chirag Sardana and 19CSU071

Semester /Section: 4th/ FS A1

Link to Code: <https://github.com/chiragsardana/Question-In-Java/tree/master/DBMS/Experiment%202>

Date:

Faculty Signature:

Remarks:

Objective

Design a Relational Database Design for the COMPANY database from the ER/EER diagram for the following set of requirements.

Program Outcome

- The students will be able to map the conceptual database design to logical (relational) database design.

Problem Statement

The COMPANY database keeps track of a company's employees, departments, and projects. Map the ER/EER diagram of the COMPANY database created in the previous experiment to Relational Database Design. Follow the below mentioned steps to successfully map ER/EER to relational tables:

- Step 1: Mapping of Regular Entity Types
- Step 2: Mapping of Weak Entity Types
- Step 3: Mapping of Binary 1:1 Relation Types
- Step 4: Mapping of Binary 1:N Relationship Types
- Step 5: Mapping of Binary M:N Relationship Types
- Step 6: Mapping of Multivalued attributes
- Step 7: Mapping of N-ary Relationship Types
- Step 8: Mapping EER Model Constructs to Relations
- Step 9: Options for Mapping Specialization or Generalization
- Step 10: Mapping of Union Types (Categories)

Background Study

- 1) Mapping of Regular Entity Types
 - For each regular (strong) entity type E in the ER schema, create a relation R that includes all the simple attributes of E. Choose one of the key attributes of E as the primary key for R. If the chosen key of E is composite, the set of simple attributes that form it will together form the primary key of R.
- 2) Mapping of Weak Entity Types

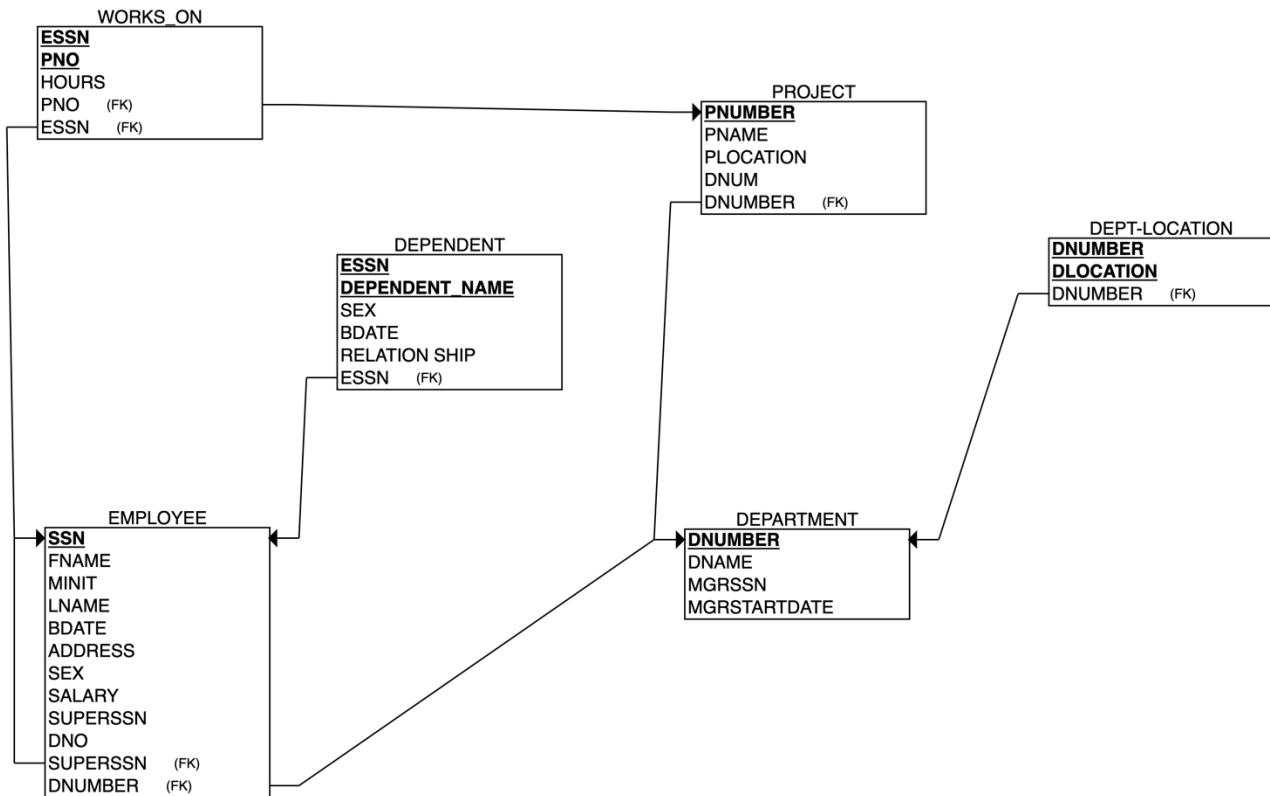
2019-20

- For each weak entity type W in the ER schema with owner entity type E, create a relation R and include all simple attributes (or simple components of composite attributes) of W as attributes of R.
 - Include as foreign key attributes of R the primary key attribute(s) of the relation(s) that correspond to the owner entity type(s).
- 3) Mapping of Binary 1:1 Relationship Types
- For each binary 1:1 relationship type R in the ER schema, identify the relations S and T that correspond to the entity types participating in R. This has 3 approaches :
 - Foreign Key Approach
 - Merged Relation Option
 - Cross-reference or Relationship Relation Option
- 4) Mapping of Binary 1:N Relationship Types
- For each regular binary 1:N relationship type R, identify the relation S that represent the participating entity type at the N-side of the relationship type.
 - Include as foreign key in S the primary key of the relation T that represents the other entity type participating in R.
- 5) Mapping of Binary M:N Relationship Types
- For each regular binary M:N relationship type R, create a new relation S to represent R.
 - Include as foreign key attributes in S the primary keys of the relations that represent the participating entity types; their combination will form the primary key of S.
- 6) Mapping of Multi-valued attributes
- For each multi-valued attribute A, create a new relation R. This relation R will include an attribute corresponding to A, plus the primary key attribute K-as a foreign key in R-of the relation that represents the entity type of relationship type that has A as an attribute.
- 7) Mapping of N-ary Relationship Types
- For each n-ary relationship type R, where $n > 2$, create a new relationship S to represent R.
 - Include as foreign key attributes in S the primary keys of the relations that represent the participating entity types.
- 8) Options for Mapping Specialization or Generalization

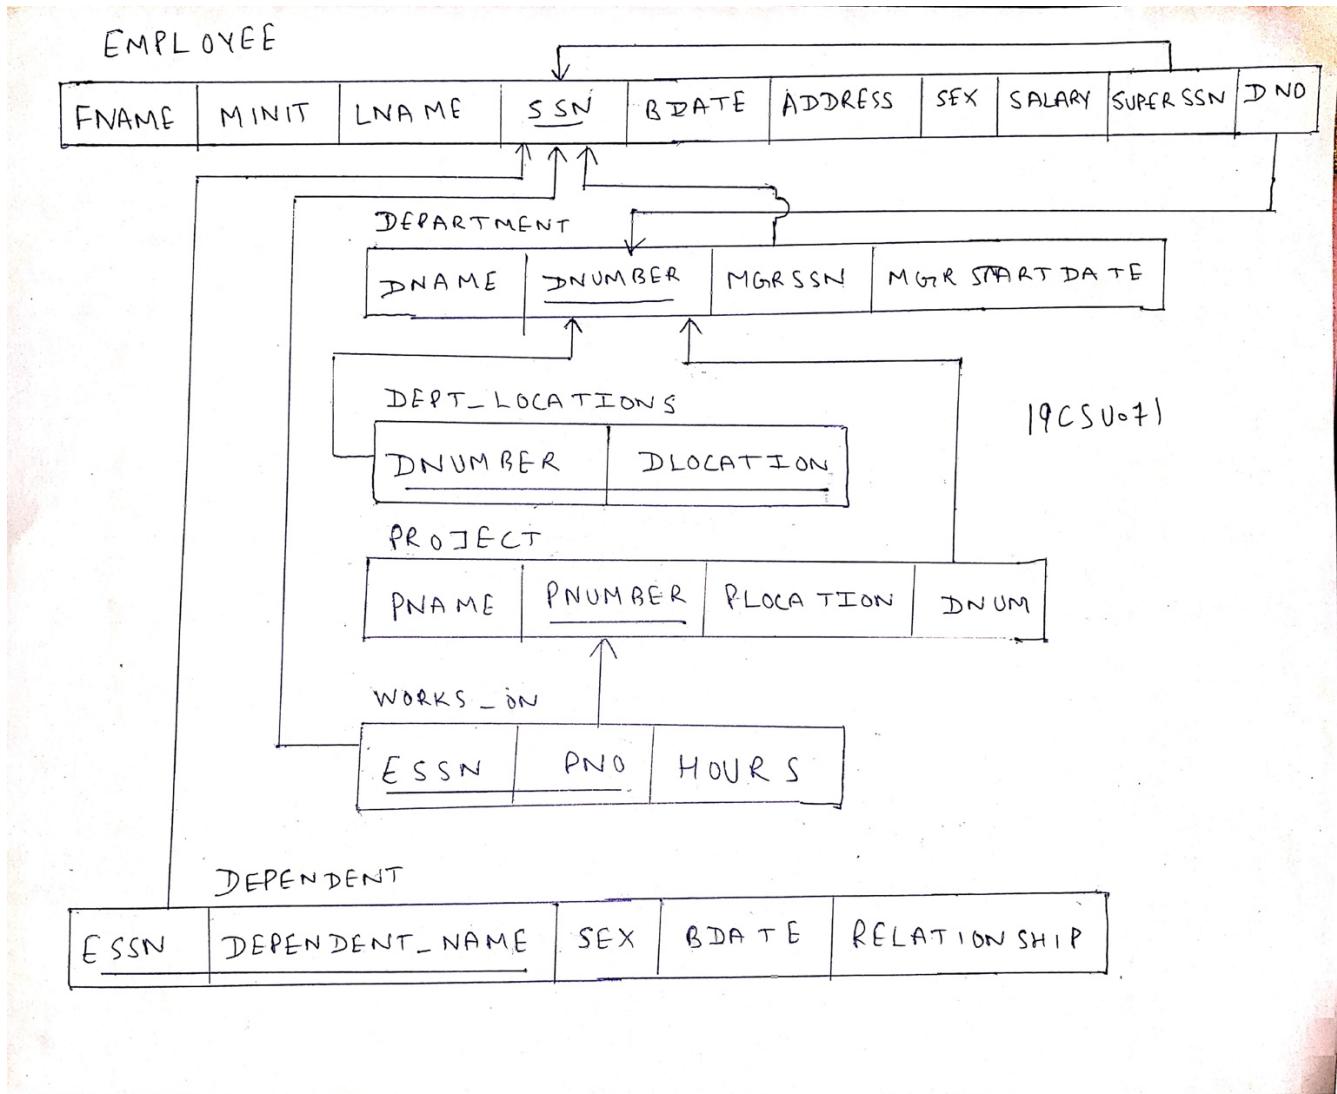
- Convert each specialization with m subclasses $\{S_1, S_2, \dots, S_m\}$ and generalized superclass C, where the attributes of C are $\{k, a_1, \dots, a_n\}$ and k is the (primary) key, into relational schemas using 1 of the following :
 - 1) Multiple relations-Superclass and subclasses.
 - 2) Multiple relations-Subclass relations only :
 - 3) Single relation with *one type attribute*
 - 4) Single relation with multiple type attributes.

- 9) Mapping of Union Types (Categories)
 - For mapping a category whose defining superclass have different keys, it is customary to specify a new key attribute, called a surrogate key, when creating a relation to correspond to the category.

Relational Database Design



2019-20



Preparatory Questions

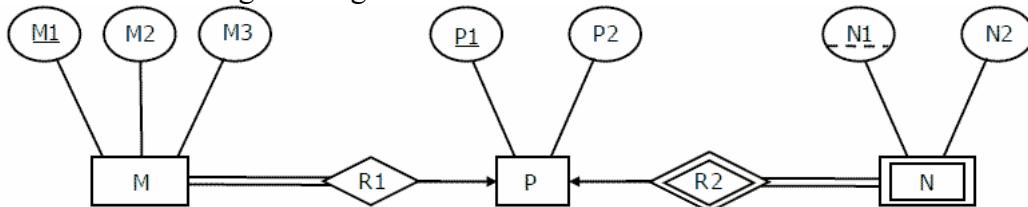
Q1) In which of the following, a separate schema is created consisting of that attribute and the primary key of the entity set.

- 1) A many-to-many relationship set
- 2) A multivalued attribute of an entity set
- 3) A one-to-many relationship set
- 4) All of the mentioned

Answer is: 2

If a multivalued dependency holds and is not implied by the corresponding functional dependency, it usually arises from this source.

Q2) Consider the following ER diagram



The minimum number of tables needed to represent M, N, P, R1, R2 is

- 1) 2
- 2) 3
- 3) 4
- 4) 5

Answer is: 2 (Means Number of Table Required is 3)

M, P are strong entities hence they must be represented by separate tables.

Many-to-one and one-to-many relationship sets that are total on the many-side can be represented by adding an extra attribute to the “many” side, containing the primary key of the “one” side. (This way no extra table will be needed for Relationship sets)

M table is modified to include primary key of P side(i.e. P1). N is weak entity, and is modified to include primary key of P (i.e, P1).

Therefore there would be minimum of 3 tables with schema given below :

M (M1, M2, M3, P1)

P (P1, P2)

N (P1, N1, N2)

Q3) Consider the data given in above question. Which of the following is a correct attribute set for one of the tables for the correct answer to the above question?

- 1) {M1, M2, M3, P1}
- 2) {M1, P1, N1, N2}
- 3) {M1, P1, N1}
- 4) {M1, P1}

Answer is: 1

Q4) Let E1 and E2 be two entities in an E/R diagram with simple single-valued attributes. R1 and R2 are two relationships between E1 and E2, where R1 is one-to-many and R2 is many-to-many. R1 and R2 do not have any attributes of their own. What is the minimum number of tables required to represent this situation in the relational model?

- 1) 2
- 2) 3
- 3) 4
- 4) 5

Answer is: 2

Strong entities E1 and E2 are represented as separate tables.

In addition to that many-to-many relationships(R2) must be converted as separate table by having primary keys of E1 and E2 as foreign keys.

One-to-many relationship (R1) must be transferred to ‘many’ side table(i.e. E2) by having primary key of one side(E1) as foreign key(this way we need not to make a separate table for R1).

Let relation schema be E1(a1,a2) and E2(b1,b2).

Relation E1(a1 is the key)

a1 a2

1 3

2019-20

2 4

3 4

Relation E2(b1 is the key, a1 is the foreign key, hence R1(one-many) relationship set satisfy here)

b1 b2 a1

7 4 2

8 7 2

9 7 3

Relation R2 ({a1, b1} combined is the key here , representing many-many relationship R2)

a1 b1

1 7

1 8

2 9

3 9

Hence we will have minimum of 3 tables.

Q5) What is the min and max number of tables required to convert an ER diagram with 2 entities and 1 relationship between them with partial participation constraints of both entities?

- 1) Min 1 and max 2
- 2) Min 1 and max 3
- 3) Min 2 and max 3
- 4) Min 2 and max 2

Answer is: 3

Maximum number of tables required is 3 in case of many to many relationships between entities. Minimum number of tables is 1 in case of unary relationship and total participation of atleast one entity. But in case of partial participation of both entities, minimum number of tables required is 2.

Q6) Which of the following has each related entity set has its own schema and there is an additional schema for the relationship set.

- 1) A many-to-many relationship set
- 2) A multivalued attribute of an entity set

2019-20

- 3) A one-to-many relationship set
- 4) All of the mentioned

Answer is: 1

If a multivalued dependency holds and is not implied by the corresponding functional dependency, it usually arises from this source.

Q7) For each attribute of a relation, there is a set of permitted values, called the _____ of that attribute.

- 1) Domain
- 2) Relation
- 3) Set
- 4) Schema

Answer is: 1

The values of the attribute should be present in the domain. Domain is a set of values permitted.

Q8) Let M and N be two entities in an E-R diagram with simple single value attributes. R1 and R2 are two relationship between M and N, where as R1 is one-to-many and R2 is many-to-many. The minimum number of tables required to represent M, N, R1 and R2 in the relational model are _____.

- 1) 4
- 2) 6
- 3) 7
- 4) 3

Answer is: 4

Strong entities E1 and E2 are represented as separate tables.

In addition to that many-to-many relationships(R2) must be converted as separate table by having primary keys of E1 and E2 as foreign keys.

One-to-many relationship (R1) must be transferred to ‘many’ side table(i.e. E2) by having primary key of one side(E1) as foreign key(this way we need not to make a separate table for R1).

Let relation schema be E1(a1,a2) and E2(b1,b2).

Relation E1(a1 is the key)

a1 a2

1 3

2 4

3 4

Relation E2(b1 is the key, a1 is the foreign key, hence R1(one-many) relationship set satisfy here)

b1 b2 a1

7 4 2

8 7 2

9 7 3

Relation R2 ({a1, b1} combined is the key here , representing many-many relationship R2)

a1 b1

1 7

1 8

2 9

3 9

Hence we will have minimum of 3 tables.

Q9) Database _____ which is the logical design of the database, and the database _____ which is a snapshot of the data in the database at a given instant in time.

- 1) Instance, Schema
- 2) Relation, Schema
- 3) Relation, Domain
- 4) Schema, Instance

Answer is: 4

Instance is an instance of time and schema is a representation.

Q10) Which option only works for a specialization whose subclasses are total (every entity in the superclass must belong to (at least) one of the subclasses) while the mapping of specialisation / generalisation form EER diagram to relational table :

- 1) Multiple relations-Superclass and subclasses.

2019-20

- 2) Multiple relations-Subclass relations only :
- 3) Single relation with *one type attribute*
- 4) Single relation with multiple type attributes.

Answer is: 2

2019-20

Experiment No: 3

Student Name and Roll Number: Chirag Sardana and 19CSU071

Semester /Section: 4th/ FS A1

Link to Code: <https://github.com/chiragsardana/Question-In-Java/tree/master/DBMS/Experiment%203>

Date:

Faculty Signature:

Remarks:

Objective

To apply SQL integrity constraints as per the DDL statements given below for COMPANY database.

Program Outcome

- The students will understand how a database is created followed by insertion of relevant data.
- The students will understand the need of applying various types of integrity constraints such as primary key, foreign key, unique key, NOT NULL, default and CHECK etc

Problem Statement:

CREATE TABLE EMPLOYEE

```
( Fname           VARCHAR(15)      NOT NULL,
  Minit          CHAR,
  Lname           VARCHAR(15)      NOT NULL,
  Ssn            CHAR(9)         NOT NULL,
  Bdate          DATE,
  Address        VARCHAR(30),
  Sex             CHAR,
  Salary          DECIMAL(10,2),
  Super_ssn      CHAR(9),
  Dno             INT              NOT NULL,
PRIMARY KEY (Ssn),
FOREIGN KEY (Super_ssn) REFERENCES EMPLOYEE(Ssn),
FOREIGN KEY (Dno) REFERENCES DEPARTMENT(Dnumber) );
```

CREATE TABLE DEPARTMENT

```
( Dname           VARCHAR(15)      NOT NULL,
  Dnumber         INT              NOT NULL,
  Mgr_ssn         CHAR(9)         NOT NULL,
  Mgr_start_date DATE,
PRIMARY KEY (Dnumber),
UNIQUE (Dname),
FOREIGN KEY (Mgr_ssn) REFERENCES EMPLOYEE(Ssn) );
```

2019-20

```

CREATE TABLE DEPT_LOCATIONS
(
    Dnumber          INT           NOT NULL,
    Dlocation        VARCHAR(15)  NOT NULL,
    PRIMARY KEY (Dnumber, Dlocation),
    FOREIGN KEY (Dnumber) REFERENCES DEPARTMENT(Dnumber) );
CREATE TABLE PROJECT
(
    Pname            VARCHAR(15)  NOT NULL,
    Pnumber          INT          NOT NULL,
    Plocation        VARCHAR(15),
    Dnum             INT          NOT NULL,
    PRIMARY KEY (Pnumber),
    UNIQUE (Pname),
    FOREIGN KEY (Dnum) REFERENCES DEPARTMENT(Dnumber) );
CREATE TABLE WORKS_ON
(
    Essn             CHAR(9)      NOT NULL,
    Pno              INT          NOT NULL,
    Hours            DECIMAL(3,1) NOT NULL,
    PRIMARY KEY (Essn, Pno),
    FOREIGN KEY (Essn) REFERENCES EMPLOYEE(Ssn),
    FOREIGN KEY (Pno) REFERENCES PROJECT(Pnumber) );
CREATE TABLE DEPENDENT
(
    Essn             CHAR(9)      NOT NULL,
    Dependent_name   VARCHAR(15)  NOT NULL,
    Sex              CHAR,
    Bdate            DATE,
    Relationship     VARCHAR(8),
    PRIMARY KEY (Essn, Dependent_name),
    FOREIGN KEY (Essn) REFERENCES EMPLOYEE(Ssn) );

```

Implement the following types of integrity constraints:

- 1) Primary Key
- 2) Foreign Key
- 3) Unique
- 4) Default
- 5) Auto-increment
- 6) Check
- 7) Not Null

Background Study

- 1) Primary Key Constraint: A column or group of columns in a table which helps us to uniquely identifies every row in that table is called a primary key. This DBMS can't be a duplicate. The same value can't appear more than once in the table.

Syntax to define a Primary key at column level:

column name datatype [CONSTRAINT constraint_name] PRIMARY KEY

Syntax to define a Primary key at table level:

[CONSTRAINT constraint_name] PRIMARY KEY (column_name1, column_name2, ...)

Rules for defining Primary key:

2019-20

- Two rows can't have the same primary key value
 - It must for every row to have a primary key value.
 - The primary key field cannot be null.
 - The value in a primary key column can never be modified or updated if any foreign key refers to that primary key.
- 2) Foreign Key (Referential integrity constraint): This constraint identifies any column referencing the PRIMARY KEY in another table. It establishes a relationship between two columns in the same table or between different tables. For a column to be defined as a Foreign Key, it should be defined as a Primary Key in the table which it is referring. One or more columns can be defined as Foreign key.

Syntax to define a Foreign key at column level:

[CONSTRAINT constraint_name] REFERENCES Referenced_Table_name(column_name)

Syntax to define a Foreign key at table level:

[CONSTRAINT constraint_name] FOREIGN KEY(column_name) REFERENCES referenced_table_name(column_name);

- 3) SQL Not Null Constraint : This constraint ensures all rows in the table contain a definite value for the column which is specified as not null. Which means a null value is not allowed.

Syntax to define a Not Null constraint:

[CONSTRAINT constraint name] NOT NULL

- 4) SQL Unique Key: This constraint ensures that a column or a group of columns in each row have a distinct value. A column(s) can have a null value but the values cannot be duplicated.

Syntax to define a Unique key at column level:

[CONSTRAINT constraint_name] UNIQUE

Syntax to define a Unique key at table level:

[CONSTRAINT constraint_name] UNIQUE(column_name)

- 5) SQL Check Constraint : This constraint defines a business rule on a column. All the rows must satisfy this rule. The constraint can be applied for a single column or a group of columns.

Syntax to define a Check constraint:

[CONSTRAINT constraint_name] CHECK (condition)

Work:

Creating Table

1. Department Table:

```
create table DEPARTMENT
(
    Dname varchar(15) not null,
    Dnumber int not null,
    Mgr_ssn char(9) not null,
    Mgr_start_date date,
    primary key(Dnumber),
    unique(Dname)
    # foreign key(Mgr_ssn) references EMPLOYEE(Ssn)
);
```

2. Employee Table:

```
create table EMPLOYEE
(
    Fname varchar(15) not null,
    Minit char,
    Lname varchar(15) not null,
    Ssn char(9) not null,
    Bdate date,
    Address varchar(30),
    Sex char,
```

2019-20

Salary decimal(10,2),
Super_ssn char(9),
Dno int not null,
primary key(Ssn),
foreign key(Super_ssn) references EMPLOYEE(Ssn),
foreign key(Dno) references DEPARTMENT(Dnumber)
);

3. Dept_Locations Table:

```
create table DEPT_LOCATIONS
(
Dnumber int not null,
Dlocation varchar(15),
PRIMARY KEY (Dnumber,Dlocation),
FOREIGN KEY (Dnumber) REFERENCES DEPARTMENT(Dnumber)
);
```

4. Project Table:

```
create TABLE PROJECT
(
Pname varchar(15) not null,
Pnumber int not null,
Plocation varchar(15),
Dnum int not null ,
primary key (Pnumber),
UNIQUE(Pname),
```

2019-20

FOREIGN KEY(Dnum) references DEPARTMENT(Dnumber)
);

5. Works_On Table:

```
create table WORKS_ON
(
    Essn char(9) not null,
    Pno int not null,
    Hours DECIMAL(3,1) not null,
    PRIMARY KEY(Essn,Pno),
    FOREIGN KEY(Essn) references EMPLOYEE(Ssn),
    FOREIGN KEY (Pno) references PROJECT(Pnumber)
);
```

6. Dependent Table:

```
CREATE table DEPENDENT
(
    Essn char(9) not null ,
    Dependent_name varchar(15) not null ,
    Sex char,
    Bdate DATE,
    Relationship varchar(8),
    PRIMARY KEY(Essn,Dependent_name),
    FOREIGN KEY(Essn) references EMPLOYEE(Ssn)
);
```

2019-20

Inserting Values in the Tables

1. Inserting Values in DepartmentTable:

```
INSERT INTO DEPARTMENT VALUES ('CSE', 1, '123456789', '2020-01-15');
```

```
INSERT INTO DEPARTMENT (Dname, Dnumber, Mgr_ssn)  
VALUES ('ECE', 2, '123456780');
```

```
INSERT INTO DEPARTMENT VALUES ('IT', 3, '123456781', '2020-10-14');
```

```
INSERT INTO DEPARTMENT (Dname, Dnumber, Mgr_ssn, Mgr_start_date)  
VALUES ('CSE-Full Stack', 4, '123456782', '2021-01-12'),  
('CSE-Cyber',5,'123456784', '2020-10-13');
```

2. Inserting Values in Employee Table:

```
INSERT INTO EMPLOYEE (Fname, Lname, Ssn, Bdate, Address, Sex, Salary, Dno)  
VALUES ('Chirag', 'Sardana', '123456782', '2000-10-13', 'Sirsa', 'M', 10000000.12,1);
```

```
INSERT INTO EMPLOYEE (Fname, Minit, Lname, Ssn, Bdate, Address, Sex, Salary, Super_ssn,  
Dno)  
VALUES ('Anuj', ' ', 'Sharma', '123456783', '2001-10-12', 'Delhi', 'M', 12000000.12,  
'123456782',1);
```

```
INSERT INTO EMPLOYEE (Fname, Lname, Ssn, Bdate, Address, Sex, Salary, Dno)  
VALUES ('Ayush', 'Singh', '123456784', '2001-05-05', 'Delhi', 'M', 12000000.12,1);
```

2019-20

INSERT INTO EMPLOYEE (Fname, Minit, Lname, Ssn, Bdate, Address, Sex, Salary, Super_ssn, Dno)

VALUES ('Deepak', ' ', 'Jindal', '123456785', '2000-08-26', 'Near DTU Delhi', 'M', 10000000.12, '123456782', 1), ('Nmit', ' ', 'Kumar', '123456786', '2000-08-26', 'Gurugram', 'M', 10000000.12, '123456785', 2);

3. Inserting Values in Dept_Locations Table:

INSERT INTO DEPT_LOCATIONS (Dnumber, Dlocation)

VALUES (1, 'Gurugram');

INSERT INTO DEPT_LOCATIONS (Dnumber, Dlocation)

VALUES (1, 'Delhi');

INSERT INTO DEPT_LOCATIONS (Dnumber, Dlocation)

VALUES (1, 'Sirsa');

4. Inserting Values in Project Table:

INSERT INTO PROJECT (Pname, Pnumber, Plocation, Dnum)

VALUES ('getRide', 71, 'Gurugram', 1);

INSERT INTO PROJECT (Pname, Pnumber, Dnum)

VALUES ('Light', 70, 2);

INSERT INTO PROJECT (Pname, Pnumber, Plocation, Dnum)

VALUES ('PortFolio', 711, 'NCU', 3);

2019-20

5. Inserting Values in Works_On Table:

```
INSERT INTO WORKS_ON (Essn, Pno, Hours)
VALUES ('123456782', 71, 12.1);
```

```
INSERT INTO WORKS_ON (Essn, Pno, Hours)
VALUES ('123456783', 711, 11.2);
```

```
INSERT INTO WORKS_ON (Essn, Pno, Hours)
VALUES ('123456784', 70, 24.2);
```

6. Inserting Values in Dependent Table:

```
INSERT INTO DEPENDENT (Essn, Dependent_name, Sex, Bdate, Relationship)
VALUES ('123456782', 'Dheeraj Sardana','M', '2000-12-10', 'Brother');
```

```
INSERT INTO DEPENDENT (Essn, Dependent_name, Sex, Bdate, Relationship)
VALUES ('123456782', 'Dheeraj Sardana','M', '2000-12-10', 'Brother');
```

```
INSERT INTO DEPENDENT (Essn, Dependent_name)
VALUES ('123456783', 'Anuj Sharma');
```

```
INSERT INTO DEPENDENT (Essn, Dependent_name, Relationship)
VALUES ('123456782', 'Deepak','Friend');
```

Output: Screenshots

**Database Management System Lab Manual
(CSL 214)**

2019-20

Command: DESC DEPARTMENT;

Output Result 1

Field	Type	Null	Key	Default	Extra
1 Dname	varchar(15)	NO	UNI	<null>	
2 Dnumber	int	NO	PRI	<null>	
3 Mgr_ssn	char(9)	NO		<null>	
4 Mgr_start_date	date	YES		<null>	

Command: DESC EMPLOYEE;

Output Result 2

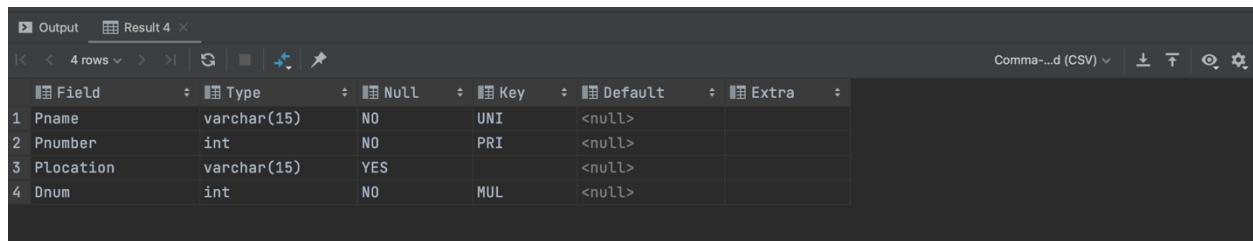
Field	Type	Null	Key	Default	Extra
1 Fname	varchar(15)	NO		<null>	
2 Minit	char(1)	YES		<null>	
3 Lname	varchar(15)	NO		<null>	
4 Ssn	char(9)	NO	PRI	<null>	
5 Bdate	date	YES		<null>	
6 Address	varchar(30)	YES		<null>	
7 Sex	char(1)	YES		<null>	
8 Salary	decimal(10,2)	YES		<null>	
9 Super_ssn	char(9)	YES	MUL	<null>	
10 Dno	int	NO	MUL	<null>	

Command: DESC DEPT_LOCATIONS;

Output Result 3

Field	Type	Null	Key	Default	Extra
1 Dnumber	int	NO	PRI	<null>	
2 Dlocation	varchar(15)	NO	PRI	<null>	

Command: DESC PROJECT;

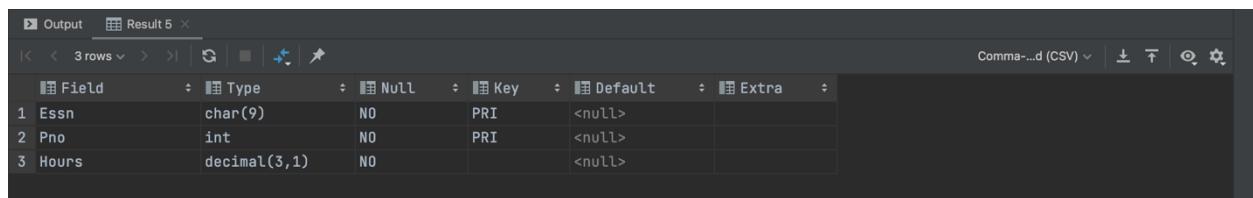


Output Result 4

Field Type Null Key Default Extra

Pname	varchar(15)	NO	UNI	<null>	
Pnumber	int	NO	PRI	<null>	
Plocation	varchar(15)	YES		<null>	
Dnum	int	NO	MUL	<null>	

Command: DESC WORKS_ON;

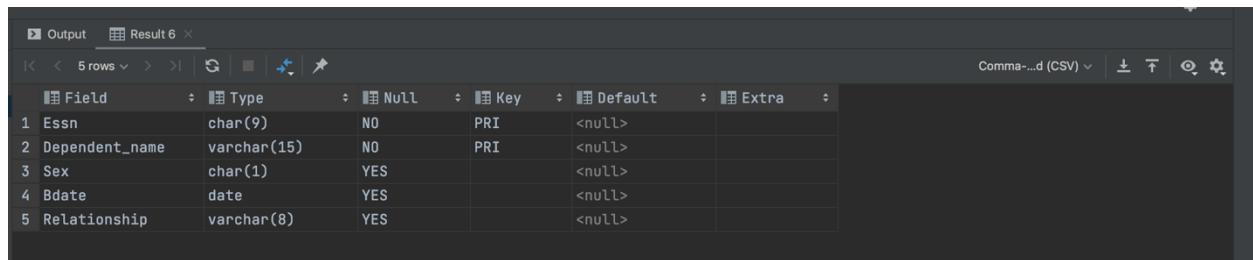


Output Result 5

Field Type Null Key Default Extra

Essn	char(9)	NO	PRI	<null>	
Pno	int	NO	PRI	<null>	
Hours	decimal(3,1)	NO		<null>	

Command: DESC DEPENDENT;

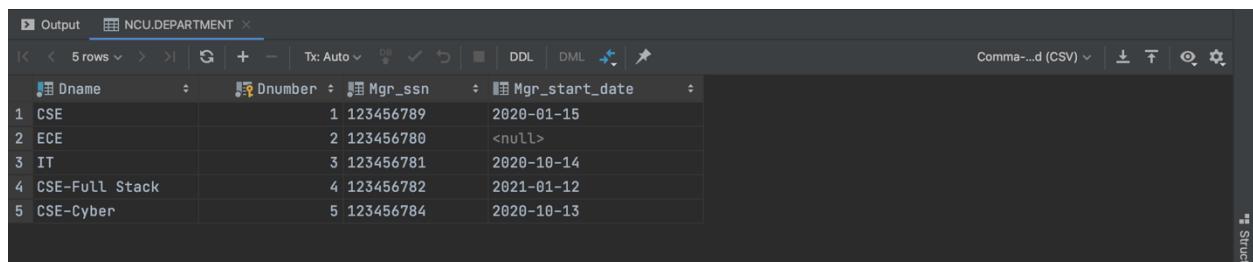


Output Result 6

Field Type Null Key Default Extra

Essn	char(9)	NO	PRI	<null>	
Dependent_name	varchar(15)	NO	PRI	<null>	
Sex	char(1)	YES		<null>	
Bdate	date	YES		<null>	
Relationship	varchar(8)	YES		<null>	

Command: SELECT * FROM DEPARTMENT;



Output NCU.DEPARTMENT

Dname Dnumber Mgr_ssn Mgr_start_date

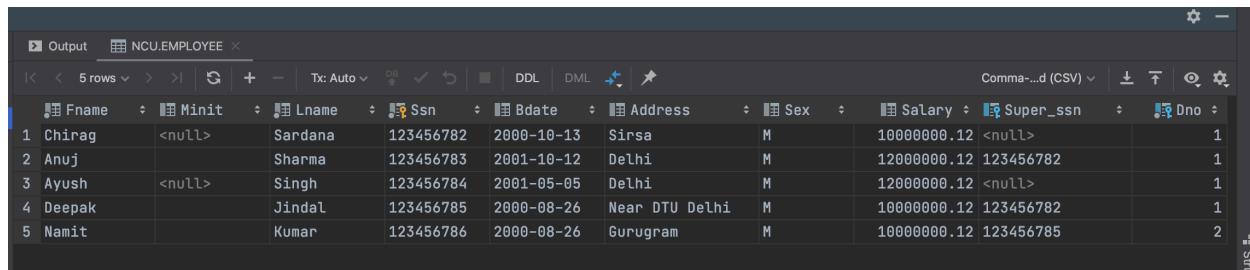
CSE	1	123456789	2020-01-15
ECE	2	123456780	<null>
IT	3	123456781	2020-10-14
CSE-Full Stack	4	123456782	2021-01-12
CSE-Cyber	5	123456784	2020-10-13

(CSL 214)

Database Management System Lab Manual

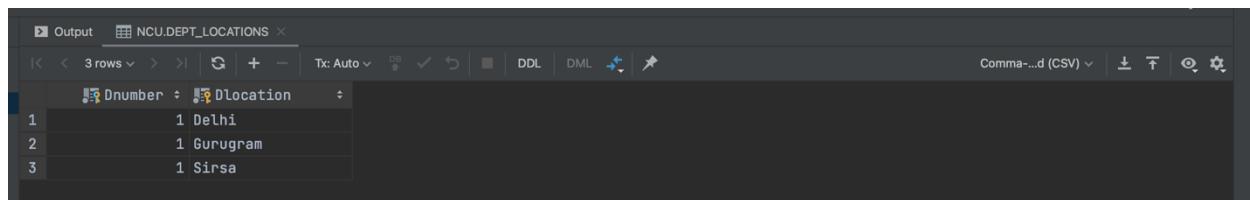
2019-20

Command: SELECT * FROM EMPLOYEE;



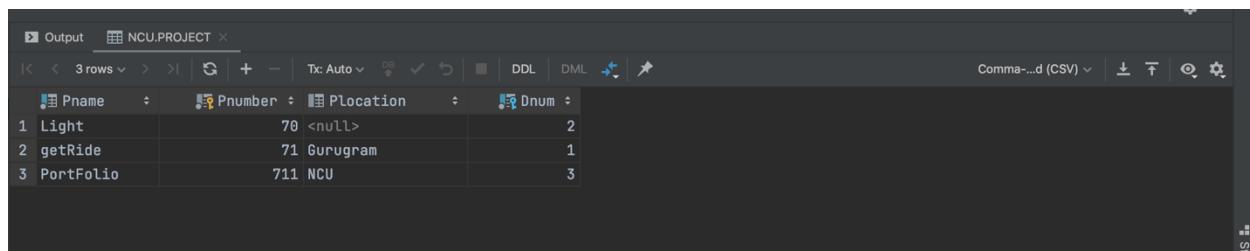
	Fname	Minit	Lname	Ssn	Bdate	Address	Sex	Salary	Super_ssn	Dno
1	Chirag	<null>	Sardana	123456782	2000-10-13	Sirsia	M	1000000.12	<null>	1
2	Anuj		Sharma	123456783	2001-10-12	Delhi	M	1200000.12	123456782	1
3	Ayush	<null>	Singh	123456784	2001-05-05	Delhi	M	1200000.12	<null>	1
4	Deepak		Jindal	123456785	2000-08-26	Near DTU Delhi	M	1000000.12	123456782	1
5	Namit		Kumar	123456786	2000-08-26	Gurugram	M	1000000.12	123456785	2

Command: SELECT * FROM DEPT_LOCATIONS;



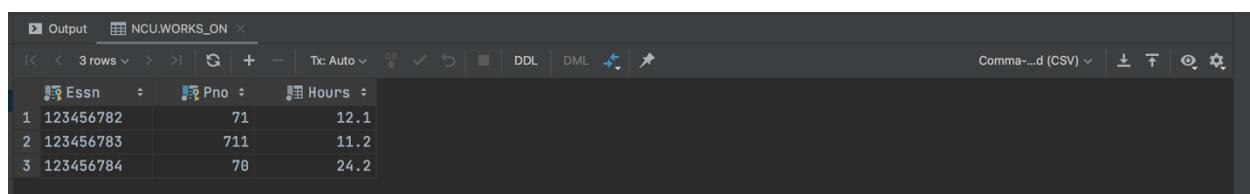
	Dnumber	Dlocation
1		1 Delhi
2		1 Gurugram
3		1 Sirsa

Command: SELECT * FROM PROJECT;



	Pname	Pnumber	Plocation	Dnum
1	Light	70	<null>	2
2	getRide	71	Gurugram	1
3	PortFolio	711	NCU	3

Command: SELECT * FROM WORKS_ON;



	Essn	Pno	Hours
1	123456782	71	12.1
2	123456783	711	11.2
3	123456784	70	24.2

2019-20

Command: SELECT * FROM DEPENDENT;

The screenshot shows a database interface with a results grid titled "NCU.DEPENDENT". The columns are labeled: Esn, Dependent_name, Sex, Bdate, and Relationship. There are four rows of data:

	Esn	Dependent_name	Sex	Bdate	Relationship
1	123456782	Chirag	M	2001-10-13	
2	123456782	Deepak	<null>	<null>	Friend
3	123456782	Dheeraj Sardana	M	2000-12-10	Brother
4	123456783	Anuj Sharma	<null>	<null>	<null>

2019-20

Preparatory Questions

- 1) Suppose (A, B) and (C, D) are two relation schemas. Let r_1 and r_2 be the corresponding relation instances. B is a foreign key that refers to C in r_2 . If data in r_1 and r_2 satisfy referential integrity constraints, which of the following is ALWAYS TRUE?
- (A) $\Pi_B(r_1) - \Pi_C(r_2) = \emptyset$
 - (B) $\Pi_C(r_2) - \Pi_B(r_1) = \emptyset$
 - (C) $\Pi_B(r_1) = \Pi_C(r_2)$
 - (D) $\Pi_B(r_1) - \Pi_C(r_2) \neq \emptyset$
- 1) A
2) B
3) C
4) D

Answer is: 1

B is a foreign key in r_1 that refers to C in r_2 . r_1 and r_2 satisfy referential integrity constraints. So every value that exists in column B of r_1 must also exist in column C of r_2 .

- 2) Given the following statements:
- S1: A foreign key declaration can always be replaced by an equivalent check assertion in SQL.
- S2: Given the table $R(a,b,c)$ where a and b together form the primary key, the following is a valid table definition.

```
CREATE TABLE S (
    a INTEGER,
    d INTEGER,
    e INTEGER,
    PRIMARY KEY (d),
    FOREIGN KEY (a) REFERENCES R)
```

Which one of the following statements is CORRECT?

- 1) S1 is TRUE and S2 is FALSE.
- 2) Both S1 and S2 are TRUE.
- 3) S1 is FALSE and S2 is TRUE.
- 4) Both S1 and S2 are FALSE

Answer is: 4

Check assertions are not sufficient to replace foreign key. Foreign key declaration may have cascade delete which is not possible by just check insertion.

Using a check condition we can have the same effect as Foreign key while adding elements to the child table. But when we delete an element from the parent table the referential integrity constraint is no longer valid. So, a check constraint cannot replace a foreign key.

So, we cannot replace it with a single check.

Foreign key in one table should uniquely identifies a row of other table. In above table definition, table S has a foreign key that refers to field ‘a’ of R. The field ‘a’ in table S doesn’t uniquely identify a row in table R.

Q3) Which of the following is not an integrity constraint?

- 1) Not null
- 2) Positive
- 3) Unique
- 4) Check ‘predicate’

Answer is: 2

Positive is a value and not a constraint.

Q4) *CREATE TABLE Manager(ID NUMERIC,Name VARCHAR(20),budget NUMERIC,Details VARCHAR(30));*

In order to ensure that the value of budget is non-negative which of the following should be used?

- 1)Check(budget>0)
- 2)Check(budget<0)
- 3)Alter(budget>0)
- 4) Alter(budget<0)

Answer is: 1

A common use of the check clause is to ensure that attribute values satisfy specified conditions, in effect creating a powerful type system.

Q5) The following table has two attributes A and C where A is the primary key and C is the foreign key referencing A with on-delete cascade.

2019-20

A	C
2	4
3	4
4	3
5	2
7	2
9	5
6	4

The set of all tuples that must be additionally deleted to preserve referential integrity when the tuple (2,4) is deleted is:

- 1) (3,4) and (6,4)
- 2) (5,2) and (7,2)
- 3) (5,2), (7,2) and (9,5)
- 4) (3,4), (4,3) and (6,4)

Answer is: 3

When (2,4) is deleted. Since C is a foreign key referring A with delete on cascade, all entries with value 2 in C must be deleted. So (5, 2) and (7, 2) are deleted. As a result of this 5 and 7 are deleted from A which causes (9, 5) to be deleted.

Q6) *CREATE TABLE course*

(...)
FOREIGN KEY (dept name) REFERENCES department. . .);

Which of the following is used to delete the entries in the referenced table when the tuple is deleted in course table?

- 1) Delete
- 2) Delete cascade
- 3) Set null
- 4) All of the mentioned

Answer is: 2

Q7) Domain constraints, functional dependency and referential integrity are special forms of ____

- 1) Foreign key
- 2) Primary key
- 3) Assertion
- 4) Referential constraint

Answer is: 3

2019-20

Q8) Which of the following can be addressed by enforcing a referential integrity constraint?

- 1) All phone numbers must include the area code
- 2) Certain fields are required (such as the email address, or phone number) before the record is accepted
- 3) Information on the customer must be known before anything can be sold to that customer
- 4) When entering an order quantity, the user must input a number and not some text (i.e., 12 rather than ‘a dozen’)

Answer is: 3

The information can be referred to and obtained.

Q9) Given relations r(w, x) and s(y, z), the result of

SELECT DISTINCT w, x

FROM r, s

is guaranteed to be same as r, provided

- 1) r has no duplicates and s is non-empty
- 2) r and s have no duplicates
- 3) s has no duplicates and r is non-empty
- 4) r and s have the same number of tuples

Answer is: 1

The query selects all attributes of r. Since we have distinct in query, result can be equal to r only if r doesn't have duplicates.

If we do not give any attribute on which we want to join two tables, then the queries like above become equivalent to Cartesian product. Cartesian product of two sets will be empty if any of the two sets is empty. So, s should have atleast one record to get all rows of r.

Q10) In SQL, relations can contain null values, and comparisons with null values are treated as unknown. Suppose all comparisons with a null value are treated as false. Which of the following pairs is not equivalent?

- 1) $x = 5 \text{ AND } \text{not}(\text{not}(x = 5))$
- 2) $x = 5 \text{ AND } x > 4 \text{ and } x < 6$, where x is an integer
- 3) $x \neq 5 \text{ AND } \text{not}(x = 5)$
- 4) None of the above

Answer is: 3

According to given question, comparison with NULL value always False, so " $x \neq 5$ " will be

2019-20

false. “x = 5” will also false and not ($x = 5$) = not(false) = true.

Hence, these are not equivalent pair.

2019-20

Experiment No: 4

Student Name and Roll Number: Chirag Sardana and 19CSU071

Semester /Section: 4th/ FS A1

Link to Code: <https://github.com/chiragsardana/Question-In-Java/tree/master/DBMS/Experiment%204>

Date:

Faculty Signature:

Remarks:

Objective

To familiarize with SELECT-FROM-WHERE SQL simple queries on the COMPANY database.

Program Outcome

- The students will be able to retrieve zero or more rows from one or more database tables or database views.

Problem Statement

From the COMPANY database as mentioned and described in the previous database :

- 1) Retrieve the birth date and address of the employee(s) whose name is ‘John B. Smith’.
- 2) Retrieve the name and address of all employees who work for the ‘Research’ department.
- 3) For every project located in ‘Stafford’, list the project number, the controlling department number, and the department manager’s last name, address, and birth date
- 4) Select all combinations of EMPLOYEE Ssn and DEPARTMENT Dname in the database.
- 5) Retrieve all the attribute values of any EMPLOYEE who works in DEPARTMENT number 5
- 6) Retrieve all distinct salary values.
- 7) Make a list of all project numbers for projects that involve an employee whose last name is ‘Smith’, either as a worker or as a manager of the department that controls the project.
- 8) Retrieve all employees whose address is in Houston, Texas.
- 9) Find all employees who were born during the 1950s
- 10) Show the resulting salaries if every employee working on the ‘ProductX’ project is given a 10 percent raise
- 11) Retrieve a list of employees and the projects they are working on, ordered by department and, within each department, ordered alphabetically by last name, then first name.

Background Study

- 1) Structured Query Language SQL contains statements for data definitions, queries, and updates (both DDL and DML)
- 2) Domain
 - o Name used with the attribute specification

2019-20

- Makes it easier to change the data type for a domain that is used by numerous attributes
- 3) Inserting values in our table using the commands :
- INSERT INTO EMPLOYEE (Fname, Lname, Dno, Ssn)
VALUES ('Richard', 'Marini', 4, '653298653')*
- 4) Out of the complete Database we create we can easily pick out and filter certain amount of data by using the SQL queries as :
- SELECT <attribute list>
FROM <table list>
WHERE <condition>;*
- <attribute list> is a list of attribute names whose values are to be retrieved by the query.
 - <table list> is a list of the relation names required to process the query.
 - <condition> is a conditional (Boolean) expression that identifies the tuples to be retrieved by the query.
- 5) Ambiguous Attribute Names : Same name can be used for two (or more) attributes as long as the attributes are in different relations, these are made to differ by mentioning their table names before the attribute names.
- 6) Aliasing, Renaming Tuple variables : We can rename the tables into some smaller easier names as per choice when it has to be used multiple times in the query.
- 7) DISTINCT keyword is used when the query wishes to derive no two duplicate values.
- 8) SET operations such as UNION and INTERSECT can also be applied on the tables to filter out the values as per choice.
- 9) Substring Pattern matching is carried out by the use of the keyword LIKE which helps in retrieving the column values of the tuple matching our mentioned substring.
- 10) ORDERBY <attributes>
- Helps in sorting the tuple values of the attributes by default set to ascending and can be changed to descending .

Work:

Before Retrieving, I am going to Insert Data More According to the Question Given Above.

*INSERT INTO EMPLOYEE (Fname, Minit, Lname, Ssn, Bdate, Address, Sex, Salary, Dno)
VALUES ('John', 'B', 'Smith', '12345677', '1990-05-05', 'Delhi', 'M', 12000000.12,1);*

2019-20

```
INSERT INTO DEPARTMENT VALUES ('Research', 6, '12345677', '2020-10-14');
```

```
INSERT INTO EMPLOYEE (Fname,Lname, Ssn, Bdate, Address, Sex, Salary, Dno)
VALUES ('Aastha', '', '12345678', '1949-05-05', 'Delhi', 'F', 12000000.12,6);
```

```
INSERT INTO PROJECT (Pname, Pnumber, Plocation, Dnum)
VALUES ('LOS', 108, 'houston', 6);
```

```
INSERT INTO EMPLOYEE (Fname,Lname, Ssn, Bdate, Address, Sex, Salary, Dno)
VALUES ('Kanu', 'Sardana', '12345679', '1950-05-05', 'Sirsa', 'M', 40000000.12,5);
```

```
INSERT INTO EMPLOYEE (Fname,Lname, Ssn, Bdate, Address, Sex, Salary, Dno)
VALUES ('Lovkesh', 'Sardana', '123456712', '1950-10-13', 'Texas', 'M', 40000000.12,5);
INSERT INTO PROJECT (Pname, Pnumber, Plocation, Dnum)
VALUES ('ProductX', 121, 'Texas', 6);
```

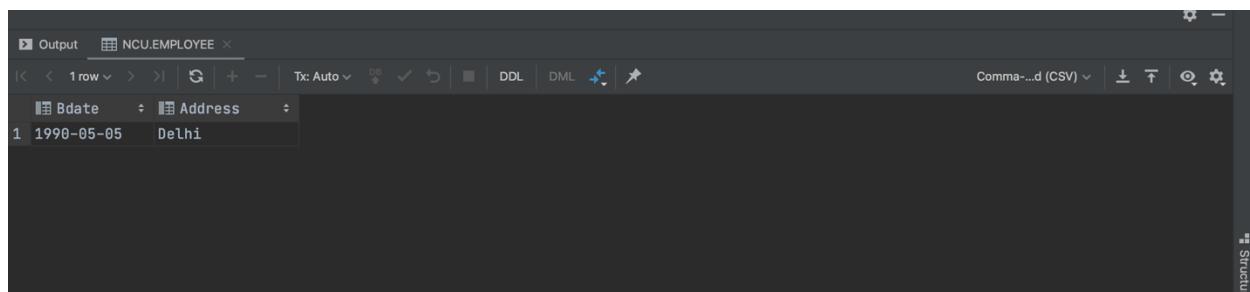
```
INSERT INTO WORKS_ON (Essn, Pno, Hours)
VALUES ('123456712', 121, 12.2);
```

Output: Screenshots

From the COMPANY database as mentioned and described in the previous database :

- 1) Retrieve the birth date and address of the employee(s) whose name is 'John B. Smith'.

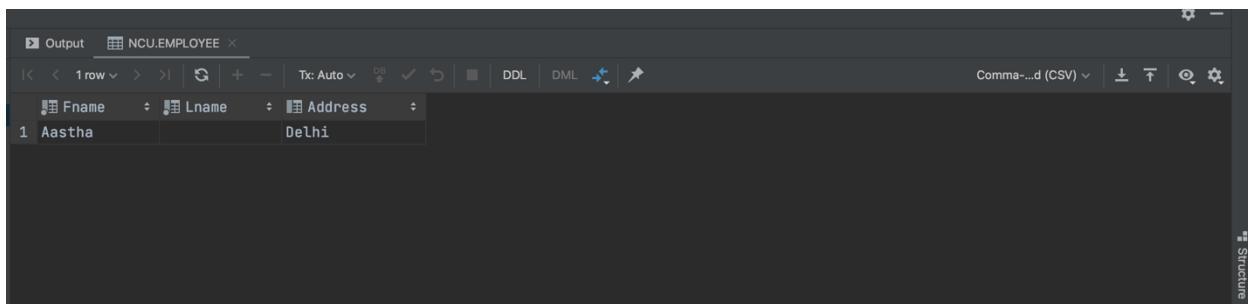
Command: SELECT Bdate, Address FROM EMPLOYEE WHERE Fname= "john" AND Minit ="B" AND Lname ="Smith";



Bdate	Address
1990-05-05	Delhi

- 2) Retrieve the name and address of all employees who work for the 'Research' department.

Command: SELECT Fname, Lname, Address FROM EMPLOYEE, Department WHERE Dname="Research" AND Dnumber=Dno;



	Fname	Lname	Address
1	Aastha		Delhi

3) For every project located in ‘houston’ , list the project number, the controlling department number, and the department manager’s last name, address, and birth date

Command: SELECT Pnumber,Dnum, Lname, Address, Bdate FROM PROJECT,DEPARTMENT,EMPLOYEE WHERE Dnum = Dnumber AND Mgr_ssn = Ssn AND Plocation="houston";



	Pnumber	Dnum	Lname	Address	Bdate
1	108	6	Simth	Delhi	1990-05-05

4) Select all combinations of EMPLOYEE Ssn and DEPARTMENT Dname in the database.

Command: SELECT Ssn, Dname FROM EMPLOYEE, Department;

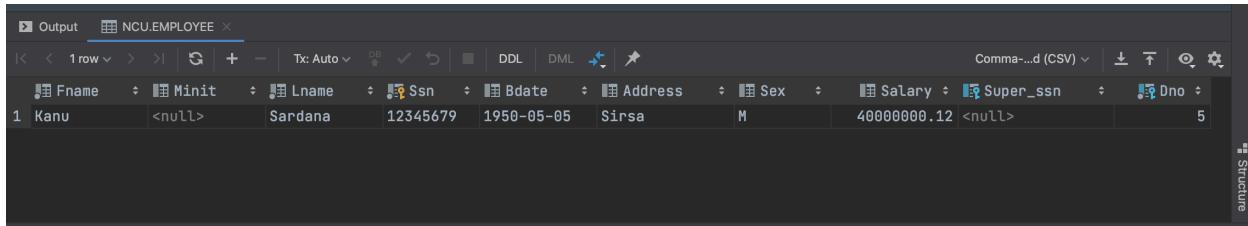
2019-20

	Ssn	Dname
1	12345677	CSE
2	12345677	CSE-Cyber
3	12345677	CSE-Full Stack
4	12345677	ECE
5	12345677	IT
6	12345677	Research
7	123456782	CSE
8	123456782	CSE-Cyber
9	123456782	CSE-Full Stack
10	123456782	ECE
11	123456782	IT
12	123456782	Research
13	123456783	CSE
14	123456783	CSE-Cyber
15	123456783	CSE-Full Stack
16	123456783	ECE
17	123456783	IT
18	123456783	Research
19	123456784	CSE
20	123456784	CSE-Cyber
21	123456784	CSE-Full Stack
22	123456784	ECE
23	123456784	IT
24	123456784	Research
25	123456785	CSE
26	123456785	CSE-Cyber
	Ssn	Dname
22	123456784	ECE
23	123456784	IT
24	123456784	Research
25	123456785	CSE
26	123456785	CSE-Cyber
27	123456785	CSE-Full Stack
28	123456785	ECE
29	123456785	IT
30	123456785	Research
31	123456786	CSE
32	123456786	CSE-Cyber
33	123456786	CSE-Full Stack
34	123456786	ECE
35	123456786	IT
36	123456786	Research
37	12345678	CSE
38	12345678	CSE-Cyber
39	12345678	CSE-Full Stack
40	12345678	ECE
41	12345678	IT
42	12345678	Research

2019-20

5) Retrieve all the attribute values of any EMPLOYEE who works in DEPARTMENT number 5

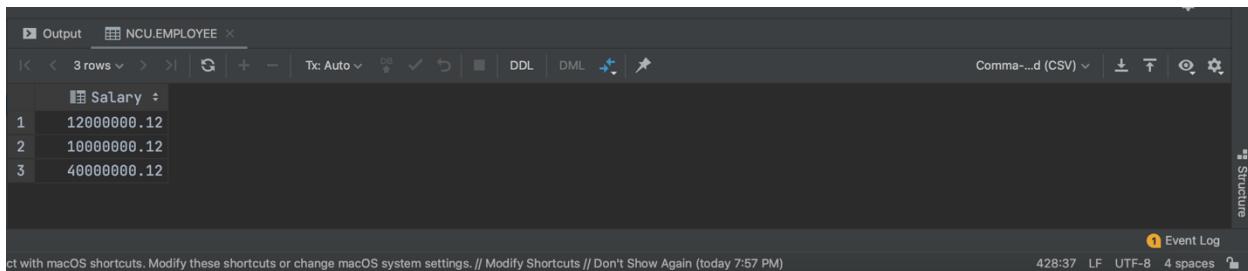
Command: SELECT * FROM EMPLOYEE WHERE Dno=5;



	Fname	Minit	Lname	Ssn	Bdate	Address	Sex	Salary	Super_ssn	Dno
1	Kanu	<null>	Sardana	12345679	1950-05-05	Sirsa	M	40000000.12	<null>	5

6) Retrieve all distinct salary values.

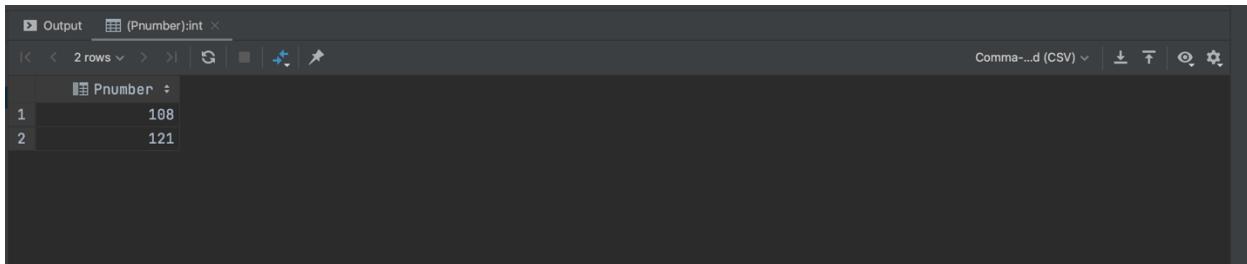
Command: SELECT DISTINCT Salary FROM EMPLOYEE;



	Salary
1	12000000.12
2	10000000.12
3	40000000.12

7) Make a list of all project numbers for projects that involve an employee whose last name is ‘Smith’, either as a worker or as a manager of the department that controls the project.

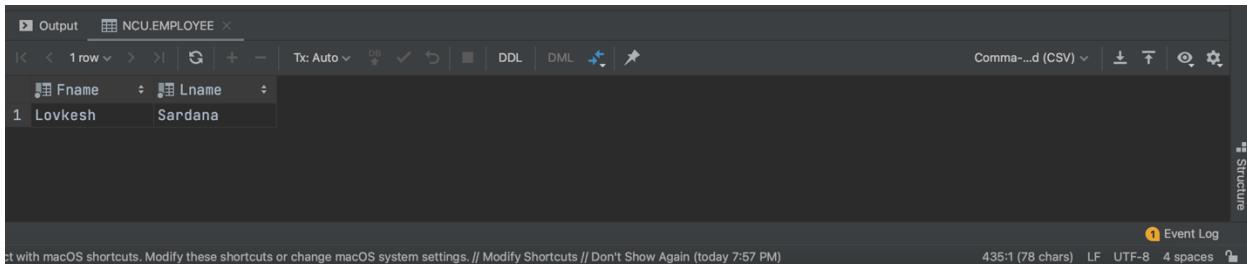
Command: (SELECT DISTINCT(Pnumber)
 FROM Project, Department, Employee
 WHERE PROJECT.Dnum = DEPARTMENT.Dnumber AND Mgr_ssn = Ssn AND
 Lname='Smith')
 UNION
 (SELECT DISTINCT(Pnumber)
 FROM Project, Works_On, Employee
 WHERE PROJECT.Pnumber = WORKS_ON.Pno AND WORKS_ON.Essn=EMPLOYEE.Ssn
 AND Lname= 'Smith')



The screenshot shows a database interface with the title bar "Output (Pnumber):int". The main area displays a table with one column labeled "Pnumber". There are two rows of data: row 1 contains the value 108, and row 2 contains the value 121. The interface includes standard navigation buttons (back, forward, search) and a toolbar with various icons.

8) Retrieve all employees whose address is in Houston OR Texas.

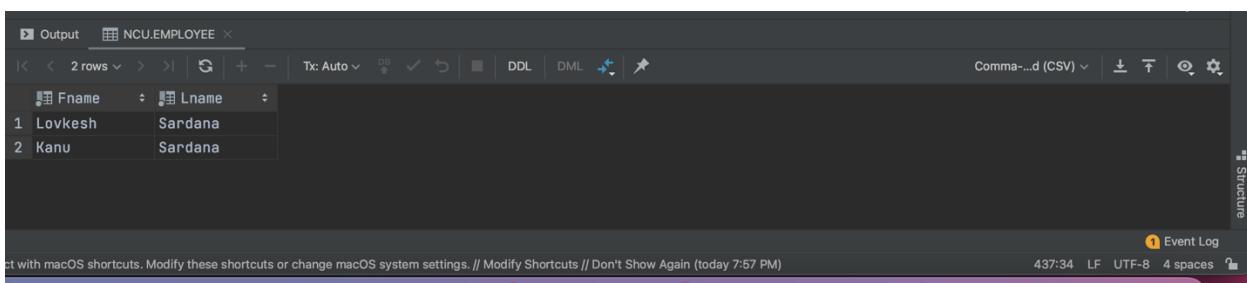
Command: SELECT Fname, Lname FROM EMPLOYEE WHERE Address LIKE "%Texas%" OR "%Houstan%";



The screenshot shows a database interface with the title bar "Output NCU.EMPLOYEE". The main area displays a table with two columns: "Fname" and "Lname". There is one row of data: row 1 contains "Lovkesh" in the Fname column and "Sardana" in the Lname column. The interface includes standard navigation buttons and a toolbar with various icons.

9) Find all employees who were born during the 1950s

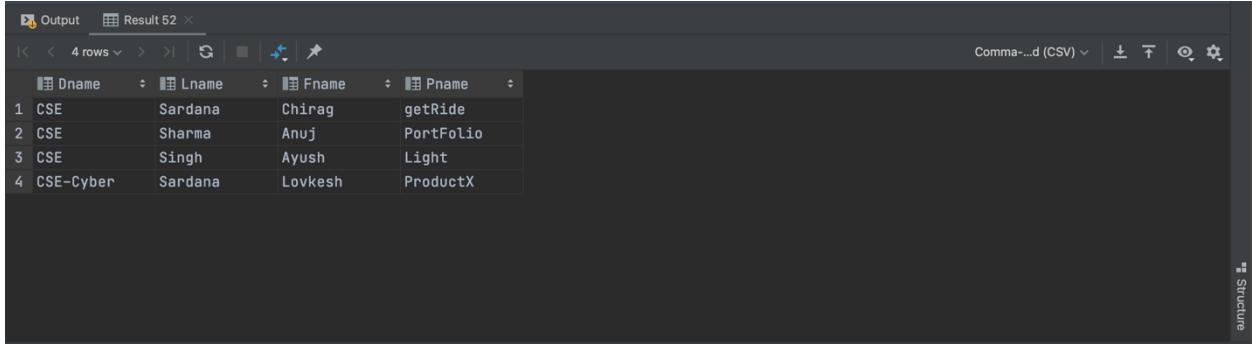
Command: SELECT Fname, Lname FROM emp WHERE Bdate LIKE "1950%";



The screenshot shows a database interface with the title bar "Output NCU.EMPLOYEE". The main area displays a table with two columns: "Fname" and "Lname". There are two rows of data: row 1 contains "Lovkesh" in the Fname column and "Sardana" in the Lname column; row 2 contains "Kanu" in the Fname column and "Sardana" in the Lname column. The interface includes standard navigation buttons and a toolbar with various icons.

10) Retrieve a list of employees and the projects they are working on, ordered by department and, within each department, ordered alphabetically by last name, then first name.

Command: SELECT D.Dname, E.Lname, E.Fname, P.Pname
 FROM Department AS D, Employee AS E, Works_On AS W, Project AS P
 WHERE D.Dnumber=E.Dno AND E.Ssn=W.Essn AND W.Pno=P.Pnumber
 ORDER BY D.Dname, E.Lname, E.Fname;

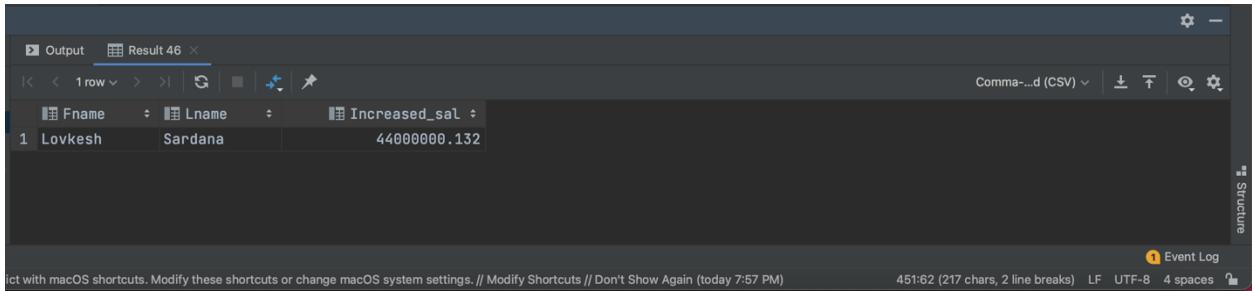


The screenshot shows a database query results window with the following table:

	Dname	Lname	Fname	Pname
1	CSE	Sardana	Chirag	getRide
2	CSE	Sharma	Anuj	PortFolio
3	CSE	Singh	Ayush	Light
4	CSE-Cyber	Sardana	Lovkesh	ProductX

11) Show the resulting salaries if every employee working on the ‘ProductX’ project is given a 10 percent raise

Command: SELECT E.Fname, E.Lname, 1.1 * E.Salary AS Increased_sal FROM EMPLOYEE AS E, WORKS_ON AS W, PROJECT AS P WHERE E.Ssn=W.Essn AND W.Pno=P.Pnumber AND P.Pname="ProductX";



The screenshot shows a database query results window with the following table:

	Fname	Lname	Increased_sal
1	Lovkesh	Sardana	44000000.132

2019-20

Preparatory Questions

Q1) Which operator performs Pattern matching ?

- 1) BETWEEN operator
- 2) LIKE operator
- 3) EXISTS operator
- 4) None of the above

Answer is: 2

LIKE is a keyword that is used in the WHERE clause. Basically, LIKE allows us to do a search based operation on a pattern rather than specifying exactly what is desired (as in IN) or spell out a range (as in BETWEEN).

Q2) In SQL which commands are used to change the storage characteristics of the table?

- 1) ALTER TABLE
- 2) MODIFY TABLE
- 3) CHANGE TABLE
- 4) All of the above

Answer is: 1

To change the structure of the table we use ALTER TABLE.

Syntax:

```
ALTER TABLE "table_name"  
ADD "column_name" datatype
```

Q3) _____ removes all rows from a table without logging the individual row deletions.

- 1) DELETE
- 2) REMOVE
- 3) DROP
- 4) TRUNCATE

Answer is: 4

TRUNCATE statement is a Data Definition Language (DDL) operation that marks the extents of a table for deallocation.

2019-20

- Q4) If you don't specify ASC or DESC after a SQL ORDER BY clause, the following is used by default :
- 1) ASC
 - 2) DESC
 - 3) There is no default value
 - 4) None of the mentioned

Answer is: 1

- Q5) What is the purpose of the SQL AS clause?

- 1) The AS SQL clause is used to change the name of a column in the result set or to assign a name to a derived column
- 2) The AS clause is used with the JOIN clause only
- 3) The AS clause defines a search condition
- 4) All of the mentioned

Answer is: 1

SQL AS clauses are defined for columns and tables to give an alias name. Basically, aliases are created to increase the readability of the query and column headings in the output.

- Q6) Which of the following command is used to delete a table in SQL?

- 1) Delete
- 2) Truncate
- 3) Remove
- 4) Drop

Answer is: 4

- Q7) Which of the following commands makes the updates performed by the transaction permanent in the database?

- 1) ROLLBACK
- 2) COMMIT
- 3) TRUNCATE
- 4) DELETE

Answer is: 2

- Q8) Select operation in SQL is equivalent to

- 1) the selection operation in relational algebra
- 2) the selection operation in relational algebra, except that select in SQL retains duplicates

2019-20

- 3) the projection operation in relational algebra
- 4) the projection operation in relational algebra, except that select in SQL retains duplicates

Answer is: 4

Select operation is equivalent to the projection operation in relational algebra, except that select in SQL retains duplicates and on the contrary projection removes the duplicates

Experiment No: 5

Student Name and Roll Number:

Semester /Section:

Link to Code:

Date:

Faculty Signature:

Remarks:

Objective

To familiarize with different JOIN operations in SQL on the COMPANY database.

Program Outcome

- The students will be to establish a connection between two or more database tables based on matching columns, thereby creating a relationship between the tables.

Problem Statement

Consider the Sample database given below and answer the queries as stated :

Pack_grades

Grade_id	Grade_name	Min_price	Max_price
----------	------------	-----------	-----------

Customers

Customer_id	First_name	Last_name	Birth_date	Join_date	City	Pack_id	State
-------------	------------	-----------	------------	-----------	------	---------	-------

Packages

Pack_id	speed	Start_date	Monthly_payment	Sector_id
---------	-------	------------	-----------------	-----------

Sectors

Sector_id	Sector_name
-----------	-------------

- Write a query to display first name, last name, package number and internet speed for all customers.
- Write a query to display first name, last name, package number and internet speed for all customers whose package number equals 22 or 27. Order the query in ascending order by last name.
- Display the package number, internet speed, monthly payment and sector name for all packages (*Packages* and *Sectors* tables).
- Display the customer name, package number, internet speed, monthly payment and sector name for all customers (*Customers*, *Packages* and *Sectors* tables).
- Display the customer name, package number, internet speed, monthly payment and sector name for all customers in the business sector (*Customers*, *Packages* and *Sectors* tables).
- Display the last name, first name, join date, package number, internet speed and sector name for all customers in the private sector who joined the company in the year 2006.

2019-20

- 7) Display the package number, internet speed, monthly payment and package grade for all packages (*Packages* and *Pack_Grades* tables).
- 8) Display the first name, last name, internet speed and monthly payment for all customers. Use INNER JOIN to solve this exercise.
- 9) Display the last name, first name and package number for all customers who have the same package number as customer named ‘Amado Taylor’ (*Customers* table).
- 10) Display the last name, first name and monthly discount for all customers whose monthly discount is lower than the monthly discount of employee number 103 (*Customers* table).
- 11) Display the package number and internet speed for all packages whose internet speed is equal to the internet speed of package number 10 (*Packages* table).

Background Study

- 1) JOINED TABLE: Permits users to specify a table resulting from a join operation in the FROM clause of a query
 - o The attributes of such a table are all the attributes of the first table followed by all the attributes of the second table.
 - o The default type of join in a joined table is called an inner join, where a tuple is included in the result only if a matching tuple exists in the other relation.
- 2) NATURAL JOIN on two relations R and S
 - o No join condition specified
 - o Implicit EQUIJOIN condition for each pair of attributes with same name from R & S
 - o It is possible to rename the attributes so that they match, if the names of the join attributes are not the same in the base relations.
- 3) LEFT OUTER JOIN
 - o Every tuple in left table must appear in result
 - o If no matching tuple, padded with NULL values for attributes of right table
- 4) RIGHT OUTER JOIN
 - o Every tuple in right table must appear in result
 - o If no matching tuple, padded with NULL values for the attributes of left table
- 5) FULL OUTER JOIN
 - o Used when both the tables are taken completely
- 6) CROSS JOIN – for Cartesian Product

Output: Screenshots

Preparatory Questions

Q1) Database table by name Loan_Records is given below.

Borrower	Bank_Manager	Loan_Amount
Ramesh	Sunderajan	10000.00
Suresh	Ramgopal	5000.00
Mahesh	Sunderajan	7000.00

What is the output of the following SQL query?

```
SELECT Count(*)  
FROM (( SELECT Borrower, Bank_Manager  
       FROM Loan_Records) AS S  
     NATURAL JOIN ( SELECT Bank_Manager, Loan_Amount  
                   FROM Loan_Records) AS T);
```

- 1) 3
- 2) 9
- 3) 5
- 4) 6

Q2) Which product is returned in a join query have no join condition:

- 1) Equijoins
- 2) Cartesian
- 3) Both Equijoins and Cartesian
- 4) None of the mentioned

Q3) Which join refers to join records from the write table that have no matching key in the left table are include in the result set:

- 1) Left outer join
- 2) Right outer join
- 3) Full outer join
- 4) Half outer join

Q4) Which operation are allowed in a join view:

- 1) UPDATE
- 2) INSERT
- 3) DELETE
- 4) All of the mentioned

Q5) Which view that contains more than one table in the top-level FROM clause of the SELECT statement:

- 1) Join view
- 2) Datable join view
- 3) Updatable join view
- 4) All of the mentioned

Q6) The following SQL is which type of join:

```
SELECT CUSTOMER_T.CUSTOMER_ID, ORDER_T.CUSTOMER_ID, NAME, ORDER_ID  
FROM CUSTOMER_T, ORDER_T
```

WHERE CUSTOMER_T.CUSTOMER_ID = ORDER_T.CUSTOMER_ID ?

- 1) Equi-join
- 2) Natural join
- 3) Outer join
- 4) Cartesian join

Q7) The following SQL is which type of join:

*SELECT CUSTOMER_T.CUSTOMER_ID, ORDER_T.CUSTOMER_ID, NAME, ORDER_ID
FROM CUSTOMER_T, ORDER_T ?*

- 1) Equi-join
- 2) Natural join
- 3) Outer join
- 4) Cartesian join

Q8) Which of the following statements is true concerning subqueries?

- 1) Involves the use of an inner and outer query
- 2) Cannot return the same result as a query that is not a subquery
- 3) Does not start with the word SELECT
- 4) All of the mentioned

Q9) Suppose ORACLE relation R(A, B) currently has tuples $\{(1, 2), (1, 3), (3, 4)\}$ and relation S(B, C) currently has $\{(2, 5), (4, 6), (7, 8)\}$. Consider the following two SQL queries SQ₁ and SQ₂:

SQ₁ : Select * From R Full Join S On R.B = S.B;

SQ₂ : Select * From R Inner Join S On R.B = S.B;

The numbers of tuples in the result of the SQL query SQ₁ and the SQL query SQ₂ are given by:

- 1) 2 and 6 respectively
- 2) 6 and 2 respectively
- 3) 2 and 4 respectively
- 4) 4 and 2 respectively

Q10)) Which of the following is true ? I. Implementation of self-join is possible in SQL with table alias. II. Outer-join operation is basic operation in relational algebra. III. Natural join and outer join operations are equivalent.

- 1) I and II are correct.
- 2) II and III are correct.
- 3) Only III is correct.
- 4) Only I is correct.

Experiment No: 6

Student Name and Roll Number:

Semester /Section:

Link to Code:

Date:

Faculty Signature:

Remarks:

Objective

To understand Aggregate functions and Group by Clause using SQL queries on the COMPANY database.

Program Outcome

- The students will understand the need of applying an aggregate function in order to get a single value from a set of values, thus, expressing the significance of the data it is computed from.

Problem Statement

Consider the COMPANY database in Experiment 1 and execute the following queries:

- Find the sum of the salaries of all employees, the maximum salary, the minimum salary, and the average salary.
- Find the sum of the salaries of all employees of the ‘Research’ department, as well as the maximum salary, the minimum salary, and the average salary in this department.
- Retrieve the total number of employees in the company.
- Retrieve the number of employees in the ‘Research’ department.
- Count the number of distinct salary values in the database
- For each department, retrieve the department number, the number of employees in the department, and their average salary.
- For each project, retrieve the project number, the project name, and the number of employees who work on that project.
- For each project *on which more than two employees work*, retrieve the project number, the project name, and the number of employees who work on the project
- For each project, retrieve the project number, the project name, and the number of employees from department 5 who work on the project.

2019-20

- 10) For each department that has more than five employees, retrieve the department number and the number of its employees who are making more than \$40,000.

Background Study

- 1) Aggregate Functions : Used to summarize information from multiple tuples into a single-tuple summary
- 2) Grouping : Create subgroups of tuples before summarizing
- 3) Built-in aggregate functions: COUNT, SUM, MAX, MIN, and AVG (NULL values discarded when aggregate functions are applied to a particular column)
- 4) Functions can be used in the SELECT clause or in a HAVING clause
- 5) Partition relation into subsets of tuples
 - o Based on grouping attribute(s)
 - o Apply function to each such group independently
- 6) GROUP BY clause: Specifies grouping attributes
- 7) If NULLs exist in grouping attribute then separate group created for all tuples with a NULL value in grouping attribute
- 8) HAVING clause provides a condition on the summary information
- 9) SELECT clause includes only the grouping attribute and the aggregate functions to be applied on each group of tuples.
- 10) WHERE clause limit the *tuples* to which functions are applied, the HAVING clause serves to choose *whole groups*

Output: Screenshots

Preparatory Questions

Q1) Which of the following statements are TRUE about an SQL query? P : An SQL query can contain a HAVING clause even if it does not have a GROUP BY clause Q : An SQL query can contain a HAVING clause only if it has a GROUP BY clause R : All attributes used in the GROUP BY clause must appear in the SELECT clause S : Not all attributes used in the GROUP BY clause need to appear in the SELECT clause

- 1) P and R
- 2) P and S
- 3) Q and R
- 4) Q and S

Q2) Consider a database table T containing two columns X and Y each of type integer. After the creation of the table, one record (X=1, Y=1) is inserted in the table. Let MX and MY denote the respective maximum values of X and Y among all records in the table at any point in time. Using MX and MY, new records are inserted in the table 128 times with X and Y values being MX+1, 2*MY+1 respectively. It may be noted that each time after the insertion, values of MX and MY change. What will be the output of the following SQL query after the steps mentioned above are carried out?

SELECT Y FROM T WHERE X=7;

- 1) 127
- 2) 255
- 3) 129
- 4) 257

Q3) The employee information in a company is stored in the relation
Employee (name, sex, salary, deptName)

Consider the following SQL query

select deptName

from Employee
where sex = 'M'
group by deptName

having avg (salary) > (select avg (salary) from Employee)

It returns the names of the department in which

- 1) the average salary is more than the average salary in the company
- 2) the average salary of male employees is more than the average salary of all male employees in the company
- 3) the average salary of male employees is more than the average salary of employees in the same department
- 4) the average salary of male employees is more than the average salary in the company

Q4) Which of the following is aggregate function in SQL?

- 1) Avg
- 2) Select
- 3) Ordered by
- 4) distinct

Q5) Observe the given SQL query and choose the correct option.

```
SELECT branch_name, COUNT (DISTINCT customer_name)
FROM depositor, account
WHERE depositor.account number = account.account number
GROUP BY branch_id
```

- 1) The query is syntactically correct but gives the wrong answer
- 2) The query is syntactically wrong
- 3) The query is syntactically correct and gives the correct answer
- 4) The query contains one or more wrongly named clauses.

Q6) Choose the correct option regarding the query

```
SELECT branch name, COUNT (DISTINCT customer name)
FROM depositor, account
WHERE depositor.account_number = account.account_number
GROUP BY branch_id
HAVING avg(balance) = 10000;
```

- 1) The having clause checks whether the query result is true or not
- 2) The having clause does not check for any condition
- 3) The having clause allows only those tuples that have average balance 10000
- 4) None of the mentioned

2019-20

Q7) Which keyword is used to rename the resulting attribute after the application of the aggregation function?

- 1) rename
- 2) as
- 3) replace
- 4) to

Q8) Consider the following ORACLE relations :

$R(A, B, C) = \{<1, 2, 3>, <1, 2, 0>, <1, 3, 1>, <6, 2, 3>, <1, 4, 2>, <3, 1, 4>\}$ $S(B, C, D) = \{<2, 3, 7>, <1, 4, 5>, <1, 2, 3>, <2, 3, 4>, <3, 1, 4>\}$.

Consider the following two SQL queries SQ_1 and SQ_2 :

$SQ_1 : \text{SELECT } R.B, \text{AVG}(S.B) \text{ FROM } R, S \text{ WHERE } R.A = S.C \text{ AND } S.D < 7 \text{ GROUP BY } R.B;$

$SQ_2 : \text{SELECT DISTINCT } S.B, \text{MIN}(S.C) \text{ FROM } S \text{ GROUP BY } S.B \text{ HAVING COUNT(DISTINCT } S.D) > 1;$

If M is the number of tuples returned by SQ_1 and N is the number of tuples returned by SQ_2 then

- 1) $M = 4, N = 2$
- 2) $M = 5, N = 3$
- 3) $M = 2, N = 2$
- 4) $M = 3, N = 3$

Q9) What values does the count(*) function ignore?

- 1) Repetitive values
- 2) Null values
- 3) Characters
- 4) Integers

Q10) State true or false: Any attribute which is present in the having clause without being aggregated must not be present in the group by clause.

- 1) True
- 2) False

Experiment No: 7

Student Name and Roll Number:

Semester /Section:

Link to Code:

Date:

Faculty Signature:

Remarks:

Objective

To familiarize with nested SQL queries on the COMPANY database.

Program Outcome

- The students will understand how to devise independent and correlated nested queries.

Problem Statement

Consider the COMPANY database in Experiment 1 and execute the following queries:

- 1) Make a list of all project numbers for projects that involve an employee whose last name is ‘Smith’, either as a worker or as a manager of the department that controls the project.
- 2) Select the Essns of all employees who work on the same project and hours as some project that employee ‘John Smith’ (whose Ssn = ‘123456789’) works on.
- 3) Return the names of employees whose salary is greater than the salary of all the employees in department 5
- 4) Retrieve the name of each employee who has a dependent with the same first name and is the same sex as the employee
- 5) Retrieve the names of employees who have no dependents
- 6) List the names of managers who have at least one dependent using EXISTS and NOT EXISTS functions
- 7) Retrieve the name of each employee who works on all the projects controlled by department number 5 using EXISTS and NOT EXISTS functions
- 8) Retrieve the names of all employees who have two or more dependents
- 9) Retrieves the names of all employees who work on only one project

Background Study

- 1) SQL allows queries that check whether an attribute value is NULL: IS or IS NOT NULL
- 2) Nested Queries: Complete select-from-where blocks within WHERE clause of another query
 - Nested Queries generally return a table (relation)
- 3) Comparison Operator IN:
 - Compares value v with a set (or multiset) of values V
 - Evaluates to TRUE if v is one of the elements in V
- 4) = ANY (or = SOME) Operator: Returns TRUE if the value v is equal to some value in the set V (equivalent to IN)
 - Other operators that can be combined with ANY (or SOME) and ALL: $>$, \geq , $<$, \leq , and \neq
- 5) ALL Operator: ($v > \text{ALL } V$) returns TRUE if the value v is greater than *all* the values in the set (or multiset) V .
- 6) Possible ambiguity among attribute names if attributes of the same name exist—one in a relation in the FROM clause of the outer query, and another in a relation in the FROM clause of the nested query.

2019-20

- Thumb Rule: reference to an unqualified attribute refers to the relation declared in the innermost nested query.
- 7) Whenever a condition in the WHERE clause of a nested query references some attribute of a relation declared in the outer query, the two queries are said to be correlated.
 - In a correlated query, the nested query is evaluated once for each tuple (or combination of tuples) in the outer query
- 8) EXISTS function: Check whether the result of a correlated nested query is empty or not.
- 9) EXISTS and NOT EXISTS are typically used in conjunction with a correlated nested query.
- 10) EXISTS(Q): returns TRUE if there is at least one tuple in the result of the nested query Q, and it returns FALSE otherwise.
- 11) NOT EXISTS(Q): returns TRUE if there are no tuples in the result of nested query Q, and it returns FALSE otherwise.

Output: Screenshots

Preparatory Questions

Q1) Database table by name Loan_Records is given below.

2019-20

Borrower	Bank_Manager	Loan_Amount
Ramesh	Sunderajan	10000.00
Suresh	Ramgopal	5000.00
Mahesh	Sunderajan	7000.00

What is the output of the following SQL query?

```
SELECT Count(*)
FROM (( SELECT Borrower, Bank_Manager
      FROM Loan_Records) AS S
  NATURAL JOIN ( SELECT Bank_Manager, Loan_Amount
      FROM Loan_Records) AS T);
```

- 1) 3
- 2) 9
- 3) 5
- 4) 6

Q2) A relational schema for a train reservation database is given below. Passenger (pid, pname, age) Reservation (pid, class, tid)

Table: Passenger

pid pname age

0	Sachin	65
1	Rahul	66
2	Sourav	67
3	Anil	69

Table : Reservation

pid class tid

0	AC	8200
1	AC	8201
2	SC	8201
5	AC	8203
1	SC	8204
3	AC	8202

What pids are returned by the following SQL query for the above instance of the tables?

```
SELECT pid
FROM Reservation ,
```

WHERE class 'AC' AND

*EXISTS (SELECT **

FROM Passenger

WHERE age > 65 AND

Passenger.pid = Reservation.pid)

- 1) 1,0
- 2) 1,2
- 3) 1,3
- 4) 1,5

Q3) Consider the following relational schema:

Suppliers(sid:integer, sname:string, city:string, street:string)

Parts(pid:integer, pname:string, color:string)

Catalog(sid:integer, pid:integer, cost:real)

Consider the following relational query on the above database:

SELECT S.sname

FROM Suppliers S

WHERE S.sid NOT IN (SELECT C.sid

FROM Catalog C

WHERE C.pid NOT IN (SELECT P.pid

FROM Parts P

WHERE P.color <> 'blue')

Assume that relations corresponding to the above schema are not empty. Which one of the following is the correct interpretation of the above query?

- 1) Find the names of all suppliers who have supplied a non-blue part.
- 2) Find the names of all suppliers who have not supplied a non-blue part.
- 3) Find the names of all suppliers who have supplied only blue parts.
- 4) Find the names of all suppliers who have not supplied only blue parts.
- 5) None

Q4) Consider the table employee(empId, name, department, salary) and the two queries Q1 ,Q2 below. Assuming that department 5 has more than one employee, and we want to find the employees who get higher salary than anyone in the department 5, which one of the statements is TRUE for any arbitrary employee table?

2019-20

Q1 : Select e.empId

From employee e

Where not exists

*(Select * From employee s where s.department = "5" and
s.salary >=e.salary)*

Q2 : Select e.empId

From employee e

Where e.salary > Any

(Select distinct salary From employee s Where s.department = "5")

- 1) Q1 is the correct query
- 2) Q2 is the correct query
- 3) Both Q1 and Q2 produce the same answer.
- 4) Neither Q1 nor Q2 is the correct query

Q5) Consider the following relational schema:

employee(empId, empName, empDept)

customer(custId, custName, salesRepId, rating)

salesRepId is a foreign key referring to empId of the employee relation. Assume that each employee makes a sale to at least one customer. What does the following query return?

SELECT empName

FROM employee E

WHERE NOT EXISTS (SELECT custId

FROM customer C

WHERE C.salesRepId = E.empId

AND C.rating <> 'GOOD');

- 1) Names of all the employees with at least one of their customers having a 'GOOD' rating.
- 2) Names of all the employees with at most one of their customers having a 'GOOD' rating.
- 3) Names of all the employees with none of their customers having a 'GOOD' rating.

4) Names of all the employees with all their customers having a ‘GOOD’ rating.

Q6) The _____ clause allows us to select only those rows in the result relation of the _____ clause that satisfy a specified predicate.

- 1) Where, from
- 2) From, select
- 3) Select, from
- 4) From, where

Q7) The relation book (title, price) contains the titles and prices of different books. Assuming that no two books have the same price, what does the following SQL query list?

select title

from book as B

where (select count()*

from book as T

where T.price > B.price) < 5

- 1) Titles of the four most expensive books
- 2) Title of the fifth most inexpensive book
- 3) Titles of the five most expensive books
- 4) Titles of the five most expensive books

Q8) Consider the following Employee table

ID salary DeptName

1	10000	EC
2	40000	EC
3	30000	CS
4	40000	ME
5	50000	ME
6	60000	ME
7	70000	CS

How many rows are there in the result of following query?

SELECT E.ID

FROM Employee E

WHERE EXISTS (SELECT E2.salary

2019-20

*FROM Employee E2
WHERE E2.DeptName = 'CS'
AND E.salary > E2.salary)*

- 1) 0
- 2) 4
- 3) 5
- 4) 6

Q9) The _____ connective tests for set membership, where the set is a collection of values produced by a select clause. The _____ connective tests for the absence of set membership.

- 1) Or, in
- 2) Not in, in
- 3) In, not in
- 4) In, or

Q10) Consider the following relation: Cinema (theatre, address, capacity)
Which of the following options will be needed at the end of the SQL query

*SELECT P1.address
FROM Cinema P1*

Such that it always finds the addresses of theatres with maximum capacity?

- 1) WHERE P1.Capacity >= All (select P2.Capacity from Cinema P2)
- 2) WHERE P1.Capacity >= Any (select P2.Capacity from Cinema P2)
- 3) WHERE P1.Capacity > All (select max(P2.Capacity) from Cinema P2)
- 4) WHERE P1.Capacity > Any (select max (P2.Capacity) from Cinema P2)

Experiment No: 8

Student Name and Roll Number:

Semester /Section:

2019-20

Link to Code:

Date:

Faculty Signature:

Remarks:

Objective

Identifying contrast between Relational Databases and NoSQL, thereby recognizing their applications.

Program Outcome

- The students will install MongoDB shell and familiarize themselves with it.

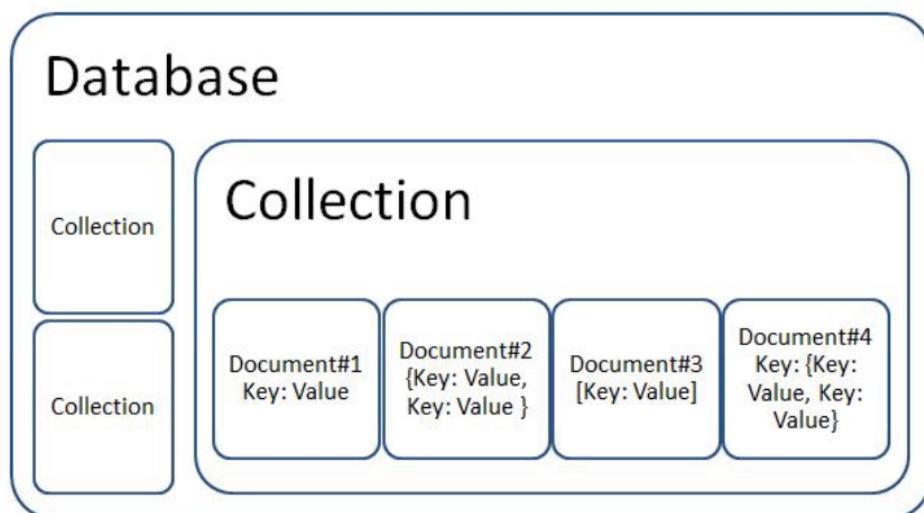
Problem Statement

MongoDB installation and shell familiarity.

Background Study

MongoDB is a Schema less database. A database in MongoDB contains collections and inside collections we have documents.

- Database is a physical container for collections. A single MongoDB server typically has multiple databases.
- Collection is a group of MongoDB documents.
- Documents within a collection can have different fields. A document is a set of key-value pairs and have dynamic schema. Dynamic schema means that documents in the same collection do not need to have the same set of fields or structure, and common fields in a collection's documents may hold different types of data.



RDBMS

- Relational database
- Need to design your tables, data structure, relations first.
- Supports SQL query language
- Table based
- Row based
- Column based
- Each row will have same number of columns
- Primary Key
- Contains schema which is predefined

MongoDB

- Non-relational and document-oriented database
- You can start coding without worrying about tables and modify your objects at a lesser cost of development.
- Supports JSON query language
- Collection based and key-value pair
- Document based
- Field based
- Each document can have different number of fields
- Primary Key (Default key **_id** provided by MongoDB itself)
- Contains dynamic schema

Output: Screenshots

Preparatory Questions

Q1) _____ command displays the list of databases.

- 1) show db
- 2) show dbs
- 3) show data
- 4) display dbs

Q2) Which of the following operation is used to switch to new database mydb ?

- 1) use dbs
- 2) use db
- 3) use mydb
- 4) use mydbs

Q3) The command to check existence of collection is :

- 1) show collection
- 2) show collections
- 3) show collect
- 4) None of the mentioned

Q4) Point out the correct statement:

- 1) A database is a set of key-value pairs
- 2) A MongoDB deployment hosts a number of databases
- 3) A document holds a set of collections
- 4) All of the mentioned

Q5) MongoDB stores all documents in:

- 1) tables
- 2) collections
- 3) rows
- 4) all of the mentioned

Q6) BSON is a binary representation of _____ documents,

- 1) JSON
- 2) XML
- 3) JScript
- 4) All of the mentioned

Q7) MongoDB documents are composed of field-and-value pairs and have the following structure:

- 1) <field1:; value1>
- 2) <field1: value1;>
- 3) <field1: value1>
- 4) none of the mentioned

Q8) The _____ field is always the first field in the document.

- 1) _id
- 2) id
- 3) Ob_id
- 4) None of the mentioned

Q9) The maximum size of a MongoDB document is

- 1) 2 MB
- 2) 16 MB
- 3) 12 MB
- 4) There is no maximum size. It depends on RAM

Q10) Which of the following statements is true?

- 1) MongoDB cannot be used as a file system.
- 2) MongoDB can run over single servers only.
- 3) Embedded documents and arrays reduce need for joins.
- 4) None

Experiment No: 9

Student Name and Roll Number:

Semester /Section:

Link to Code:

Date:

Faculty Signature:

Remarks:

Objective

Create COMPANY database using NoSQL database - MongoDB.

Program Outcome

- The students will be able to perform basic CRUD operations in MongoDB.

Problem Statement

- 1) Create a COMPANY database with the following collections:
 - Employee with fields Emp ID, Ename, Age, Mobile, Email, Address, Dno
 - Department with fields Dnumber, Dname, Dlocation
 - Project with fields Pname, Pnumber, Plocation, ControlDept
- 2) Insert 5 documents in each collection using different versions of “Insert” command.
- 3) Select all documents in a collection.
- 4) Retrieve the details of all projects that are being run in department number 50.
- 5) Retrieve the details of the first project that is being run in department number 50 in formatted manner.
- 6) Return the formatted/structured project details of all projects in department number “5” and project name as “ProductZ”.
- 7) Return the project details (including project name, project number and department number and excluding project location and default primary key) of all projects with department number “5” or project name as “ProductZ”.
- 8) Return the formatted/structured project details of all projects with department number less than “10” and project name as “ProductZ” or project location as “Delhi”.

Background Study

- 1) **use:** This command is used to create a database in MongoDB. It will return an existing database or will create a new database if it doesn't exist.
- 2) **createCollection():** To create a collections in a database
- 3) MongoDB provides the following methods to insert documents into a collection:
- 4) **db.collection.insert():** used to insert one or multiple documents. To insert multiple documents in a single query, pass an array of documents in insert() command.
- 5) **db.collection.insertOne():** used to insert only one document.
- 6) **db.collection.insertMany():** used to insert multiple documents. To insert multiple documents in a single query, pass an array of documents in insert() command.
- 7) **db.collection.find():** this method is provided in MongoDB to read documents from a collection. It returns a **cursor** to the matching documents. db.collection.find(<query>

filter>, <projection>) accepts second optional parameter that is list of fields that you want to retrieve in the form { field1: <value>, field2: <value> ... }.

- 8) **db.collection.findOne()**: it returns the first occurrence of document, otherwise null.

Output: Screenshots

Preparatory Questions

Q1) Which of the following method is used to query documents in collections ?

- 1) find
- 2) move
- 3) shell
- 4) replace

Q2) When you query a collection, MongoDB returns a _____ object that contains the results of the query.

- 1) row
- 2) cursor

2019-20

- 3) columns
- 4) none of the mentioned

Q3) Which of the following method is called while accessing documents using the array index notation ?

- 1) cur.toArray()
- 2) cursor.toArray()
- 3) doc.toArray()
- 4) all of the mentioned

Q4) The mongo shell and the drivers provide several cursor methods that call on the cursor returned by the _____ method to modify its behavior.

- 1) cursor()
- 2) find()
- 3) findc()
- 4) none of the mentioned

Q5) Which of the following method corresponds to Order by clause in SQL ?

- 1) sort()
- 2) order()
- 3) orderby()
- 4) all of the mentioned

Q6) A query may include a _____ that specifies the fields from the matching documents to return.

- 1) selection
- 2) projection
- 3) union
- 4) none of the mentioned

Q7) In MongoDB, _____ operations modify the data of a single collection.

- 1) CRUD
- 2) GRID
- 3) READ
- 4) All of the mentioned

Q8) Which of the following operation adds a new document to the users collection ?

- 1) add
- 2) insert
- 3) truncate
- 4) drop

Q9) Which of the following is a valid MongoDB JSON document:

- 1) {}
- 2) { "user_id":1,
 "user_name":"Joe Sanders",
 "occupation":["engineer","writer"] }
- 3) { "user_id":1; }

2019-20

```

    "user_name":"Joe Sanders";
    "occupation":["engineer","writer"]
}
4) {
    "user_id":1,
    "user_name":"Joe Sanders",
    "occupation":[ "occupation1":"engineer", "occupation2":"writer"]
}

```

Q10) Which of the following commands will return all the posts with number of likes greater than 100 and less than 200, both inclusive?

- 1) db.posts.find({ likes : { \$gt : 100, \$lt : 200 } });
- 2) db.posts.find({ likes : { \$gte : 100, \$lt : 200 } });
- 3) db.posts.find({ likes : { \$gt : 100 , \$lte : 200 } });
- 4) db.posts.find({ likes : { \$gte : 100 , \$lte : 200 } });

Experiment No: 10

Student Name and Roll Number:

Semester /Section:

Link to Code:

Date:

Faculty Signature:

Remarks:

Objective

Retrieve data from NoSQL database - MongoDB.

Program Outcome

- The students will be able to perform advanced CRUD operations in MongoDB.

Problem Statement

Create a database with the following collections:

Project

Possible Fields:

- Project Name (Pname)
- Project Number (Pnumber)
- Department Number in which project is being run (Dnum)
- Project Locations (Plocation)
 - Plocation can be an array of multiple locations (or)
 - Plocation can be a composite field with sub-fields – city, state, country

Department

Possible Fields:

- Department Number (Dno)
- Department Name (Dname)
- Department Manager (Dmanager)

Execute the following queries:

- 1) Insert multiple documents (2-3) together in the Project collection with atleast 1 project location as a composite attribute
- 2) Write the “Select * from Project” equivalent query in MongoDB.
- 3) Write a MongoDB query to update the project location of all projects to “Pune” and department number to “20” with project name “ProjectX”.
- 4) Write a MongoDB query to delete all the documents of collection “Project” with department number “4”.
- 5) Write a MongoDB query to retrieve the project details of all projects with project location as “Delhi” and “Mumbai” both assuming the following documents exist.

```
{ "_id" : ObjectId("5cc7fe36e37496f1a5e7e268"), "Pname" : "ProductK", "Pnumber" : 45, "Plocation" : ["Delhi", "Mumbai", "Pune"], "Dnum" : 3 }
{ "_id" : ObjectId("5cc7fe36e37496f1a5e7e268"), "Pname" : "ProductT", "Pnumber" : 56, "Plocation" : ["Delhi", "Mumbai"], "Dnum" : 15 }
{ "_id" : ObjectId("5cc7fe36e37496f1a5e7e268"), "Pname" : "ProductU", "Pnumber" : 5, "Plocation" : "Delhi", "Dnum" : 6 }
{ "_id" : ObjectId("5cc7fe36e37496f1a5e7e268"), "Pname" : "ProductL", "Pnumber" : 4, "Plocation" : "Mumbai", "Dnum" : 7 }
```

- 6) Write a MongoDB query creating a One to Many Relation between a document of “Project” Collection and “Department” Collection.
- 7) Write a MongoDB query to retrieve the project details → project name, department number, department name and department manager of projects with department manager (for department in which the project is being run) as “ABC”.

Background Study

- 1) `pretty()` is used to display the query obtained in a neat manner with all the fields clearly depicted one below the other.
- 2) Data for a particular field can also be stored as an array with multiple values for the field.
- 3) `$all` is used to match the array items for the particular field.
 - o `db.project.find({Plocation:{$all:["Delhi","Mumbai"]}}).pretty()`
- 4) `$elemMatch`: This operator matches documents that contain an array field with at least one element that matches all the specified query criteria.
 - o Syntax: { <field>: { \$elemMatch: { <query1>, <query2>, ... } } }
- 5) Fields can be embedded in the field too and while using find in such a case, always refer the field embedded in another field by using a dot notation and in quotes.
- 6) AND/OR Operations

Operation	Syntax	Example
AND	<code>>db.collection_name.find({\$and: [{key1: value1}, {key2:value2}]}))</code>	<code>>db.project.find({\$and:[{Dnum:5}, {Pname:"ProductZ"}]}).pretty();</code>
OR	<code>>db.collection_name.find({\$or: [{key1: value1}, {key2:value2}]}))</code>	<code>>db.project.find({\$or:[{Dnum:5}, {Pname:"ProductZ"}]}).pretty();</code>

- 7) Documents can either be embedded in one to one relation or one to many relation
 - o In one to one, the new field contains only a single embedded document.
 - o In one to many, the new field contains an array of multiple embedded documents.

Output: Screenshots

Preparatory Questions

Q1) The order of documents returned by a query is not defined unless you specify a _____

- 1) sortfind()
- 2) sortelse()
- 3) sort()
- 4) none of the mentioned

Q2) In aggregation pipeline, the _____ pipeline stage provides access to MongoDB queries.

- 1) \$catch
- 2) \$match
- 3) \$batch
- 4) All of the mentioned

Q3) Point out the wrong statement.

- 1) sort() modifier sorts the results by age in ascending order
- 2) Queries in MongoDB return all fields in all matching documents by default
- 3) To scale the amount of data that MongoDB sends to applications, include a projection in the queries.
- 4) None of the mentioned

2019-20

Q4) Which of the following query selects documents in the records collection that match the condition { "user_id": { \$lt: 42 } }?

- 1) db.records.findOne({ "user_id": { \$lt: 42 } }, { "history": 0 })
- 2) db.records.find({ "user_id": { \$lt: 42 } }, { "history": 0 })
- 3) db.records.findOne({ "user_id": { \$lt: 42 } }, { "history": 1 })
- 4) db.records.select({ "user_id": { \$lt: 42 } }, { "history": 0 })

Q5) Which of the following is not a projection operator?

- 1) \$slice
- 2) \$elemMatch
- 3) \$
- 4) None of the mentioned

Q6) _____ is a client or database-generated identifier that uniquely identifies this message.

- 1) messageLength
- 2) responseTo
- 3) requestID
- 4) all of the mentioned

Q7) When inserts, updates and deletes have a _____ write concern, write operations return quickly.

- 1) strong
- 2) weak
- 3) average
- 4) very strong

Q8) Which of the following MongoDB query is equivalent to the following SQL query:

UPDATE users SET status = "C" WHERE age > 25

1)

```
db.users.update(  
  { age: { $gt: 25 } },  
  { status: "C" })
```

2)

```
db.users.update(  
  { age: { $gt: 25 } },  
  { $set: { status: "C" } })
```

3)

```
db.users.update(  
  { age: { $gt: 25 } },  
  { $set: { status: "C" } },  
  { multi: true })
```

4)

```
db.users.update(  
  { age: { $gt: 25 } },  
  { status: "C" },  
  { multi: true })
```

Q9) Consider that you have a collection called population which has fields state and city. Which of the following query will calculate the population grouped by state and city?

- 1) db.population.aggregate([{ \$group: { _id: { state: "\$state", city: "\$city" }, pop: { \$sum: "\$pop" } } }])
- 2) db.population.aggregate([{ \$group: { _id: { state: "\$state", city: "\$city" }, pop: { \$sum: 1 } } }])
- 3) db.population.aggregate([{ \$group: { _id: { state: "\$state", city: "\$city" }, pop: { "\$pop": 1 } } }])
- 4) db.population.aggregate([{ \$group: { _id: { city: "\$city" }, pop: { \$sum: "\$pop" } } }])

Q10) Which of the following command is used to get all the indexes on a collection?

- 1) db.collection.getIndexes()
- 2) db.collection.showIndexes()
- 3) db.collection.findIndexes()
- 4) db.showIndexes()

Database Management System Lab Manual

2019-20